

# **Machine Learning (and Deep Learning) in just over 200 minutes with R**

## **From basic linear models to neural networks**

**Presented at the 2021 SAGI Symposium via Zoom**

**Presented by Yoni Nazarathy**

<https://yoninazarathy.com/>

**R Material created by Ajay Hemanth**

<https://www.linkedin.com/in/ajayhemanth/>

These slides are available at: <https://yoninazarathy.com/TBD>  
R source code is available at: <https://github.com/TBD>



# On the menu

1. The state of the art and the ML world
2. The ML workflow and common “practice” data sets
3. Some tools you know - used for ML
4. A walk in the random forest
5. Going deep

Who are you?

Biometricalians, using R, using specialized software.  
Strong understanding of statistics.

- A: Linear Classifiers.
- B: Google TensorFlow Ground
- C: Dense Neural Nets (MLP)
- D: Random Forests
- E: Convolutional Neural Nets

# About the speaker

<https://yoninazarathy.com/>

I am an Associate Professor at the [School of Mathematics and Physics](#) of The University of Queensland.

My expertise is in [Machine Learning](#), [Applied Probability](#), [Statistics](#), [Operations Research](#), [Simulation](#), [Scientific Computing](#), [Control Theory](#), [Queueing Theory](#), [Scheduling](#), and [Mathematical Education](#).

Application areas of my work include [epidemics](#), [wireless communication networks](#), [bio-statistics](#), [agriculture](#), [power systems](#), [software and hardware design](#), [manufacturing logistics](#), [healthcare logistics](#), and [road traffic networks](#).

## Current “Main” Project: Safe Blues

<https://safeblues.org/>

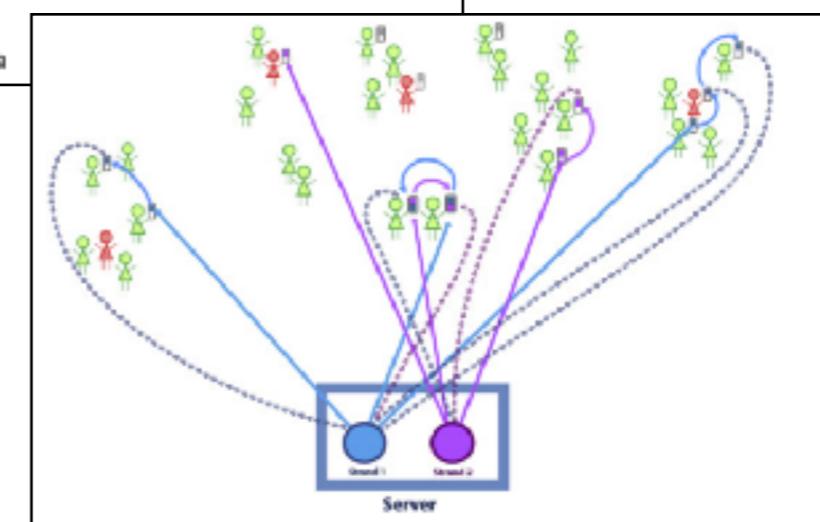
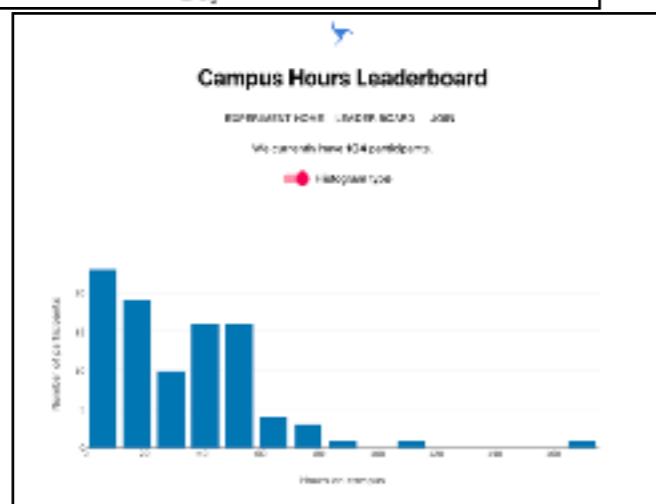
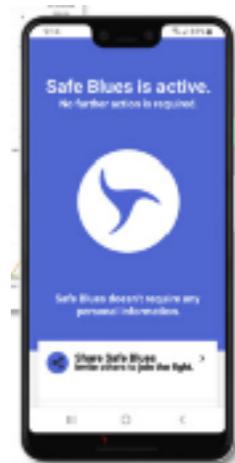
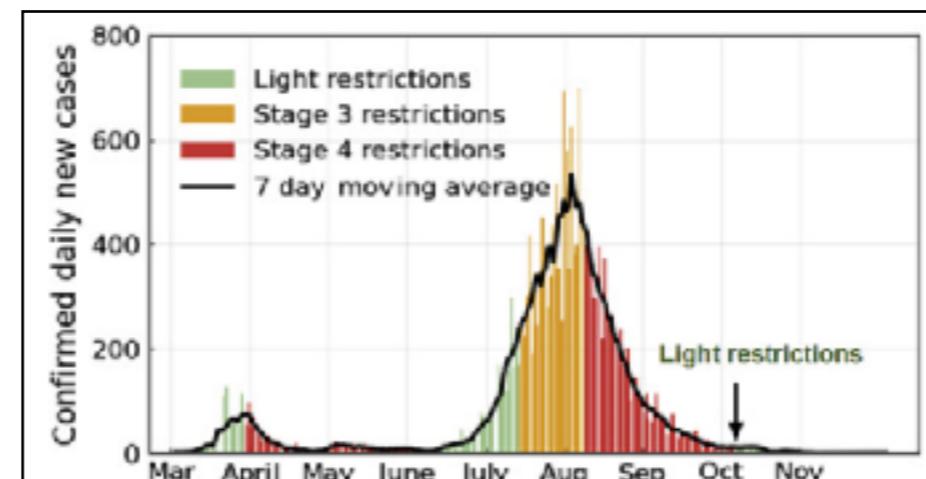
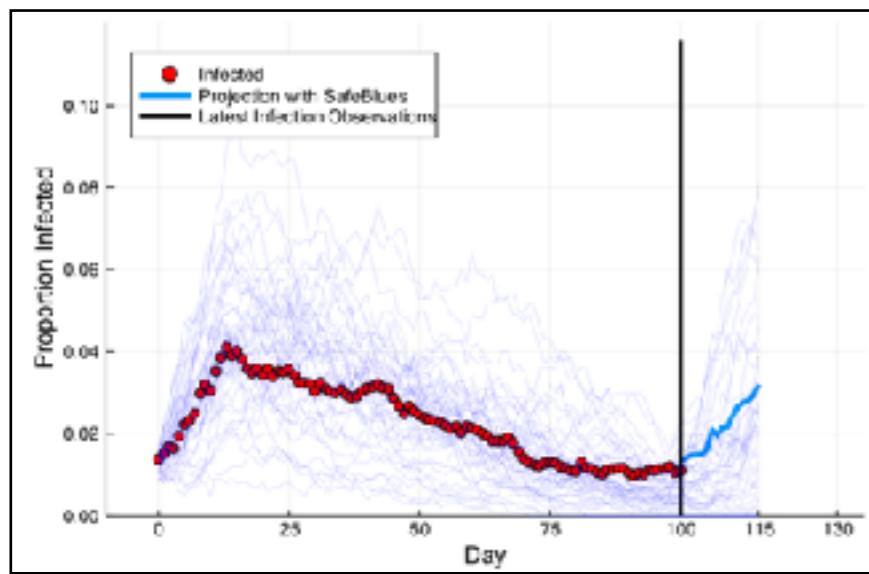
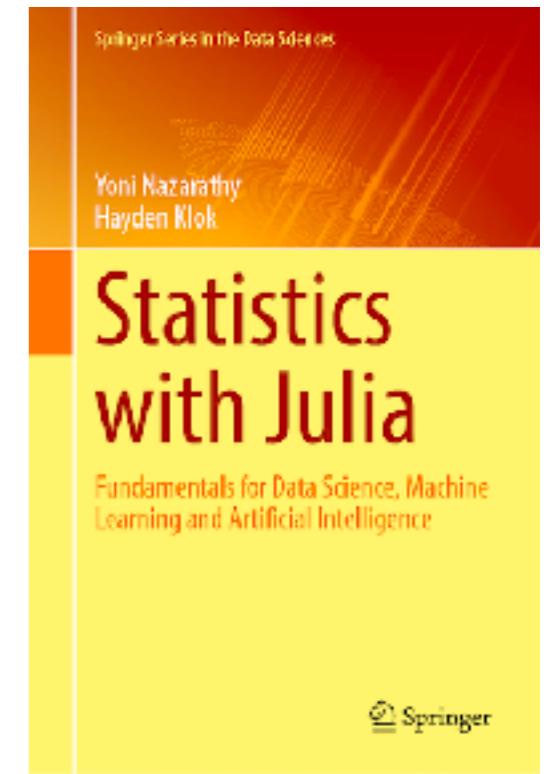
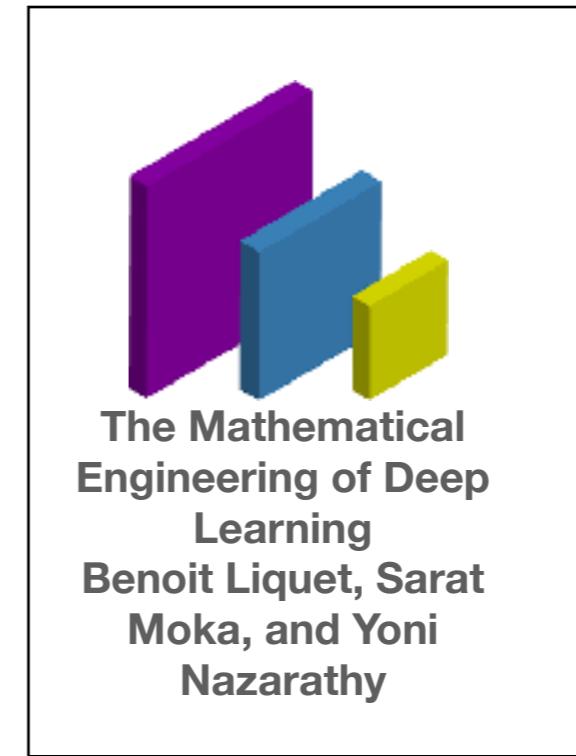
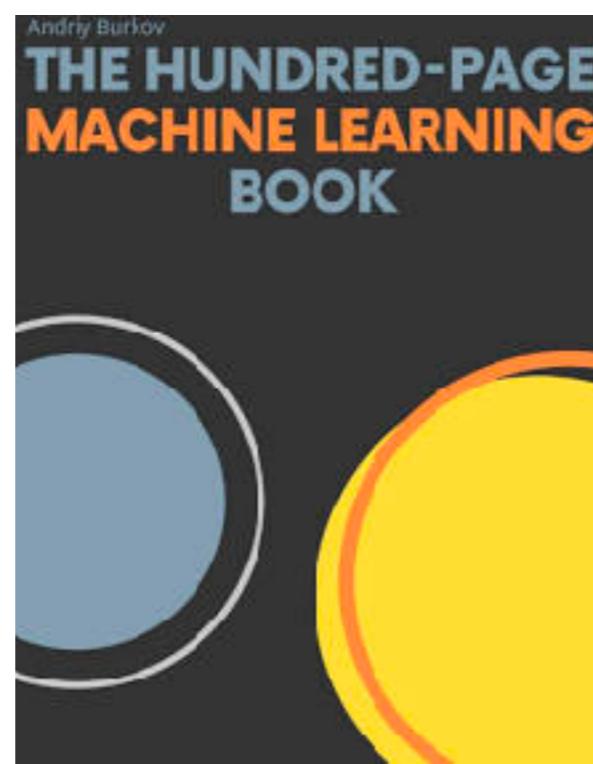


Figure 1 The 2020 outbreak in Victoria

# Some resources



<https://people.smp.uq.edu.au/DirkKroese/DSML/>

<http://themlbook.com/>

<https://deeplearningmath.org/>

<https://statisticswithjulia.org/>

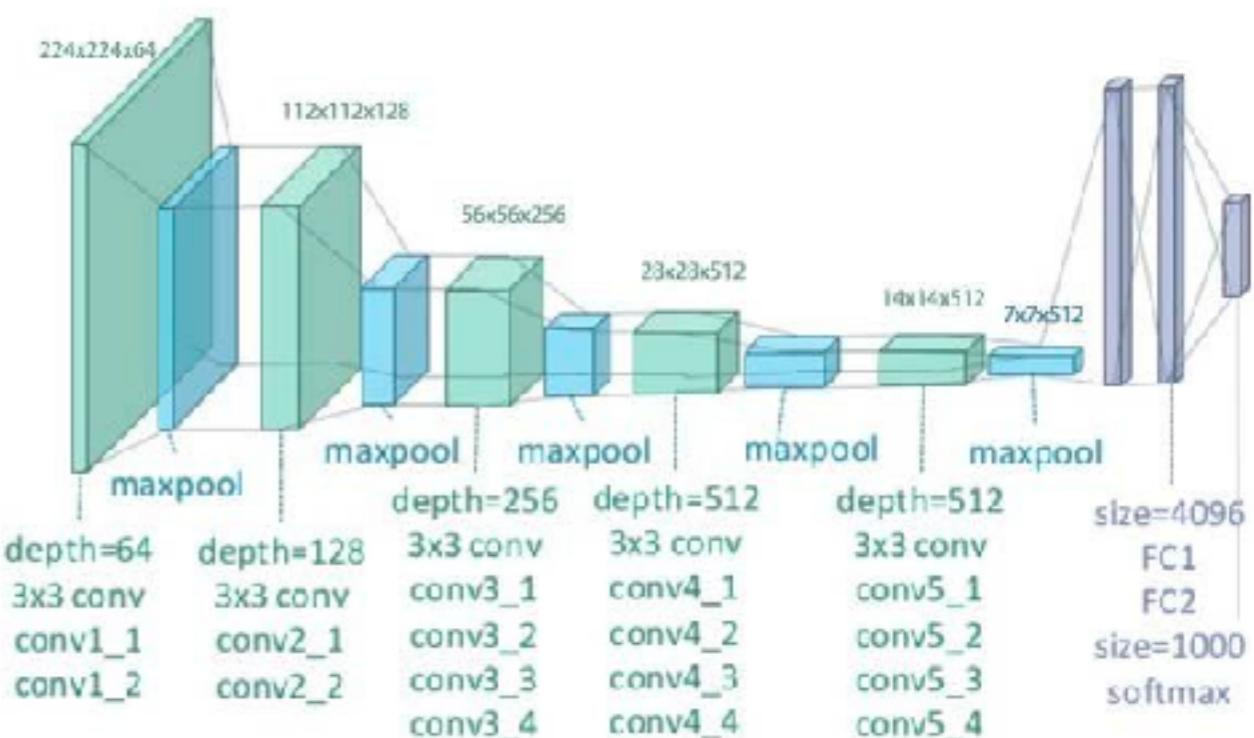
# **Part 1: The state of the art and the ML world**

```

using Metalhead
vgg = VGG19() #downloads about 0.5Gb of a pretrained neural network from the web
for (i,l) in enumerate(vgg.layers)
    println(i,": ", l)
end

1: Conv((3, 3), 3=>64, relu)
2: Conv((3, 3), 64=>64, relu)
3: MaxPool((2, 2), pad = (0, 0, 0, 0), stride = (2, 2))
4: Conv((3, 3), 64=>128, relu)
5: Conv((3, 3), 128=>128, relu)
6: MaxPool((2, 2), pad = (0, 0, 0, 0), stride = (2, 2))
7: Conv((3, 3), 128=>256, relu)
8: Conv((3, 3), 256=>256, relu)
9: Conv((3, 3), 256=>256, relu)
10: Conv((3, 3), 256=>256, relu)
11: MaxPool((2, 2), pad = (0, 0, 0, 0), stride = (2, 2))
12: Conv((3, 3), 256=>512, relu)
13: Conv((3, 3), 512=>512, relu)
14: Conv((3, 3), 512=>512, relu)
15: Conv((3, 3), 512=>512, relu)
16: MaxPool((2, 2), pad = (0, 0, 0, 0), stride = (2, 2))
17: Conv((3, 3), 512=>512, relu)
18: Conv((3, 3), 512=>512, relu)
19: Conv((3, 3), 512=>512, relu)
20: Conv((3, 3), 512=>512, relu)
21: MaxPool((2, 2), pad = (0, 0, 0, 0), stride = (2, 2))
22: #44
23: Dense(25088, 4096, relu)
24: Dropout(0.5)
25: Dense(4096, 4096, relu)
26: Dropout(0.5)
27: Dense(4096, 1000)
28: softmax

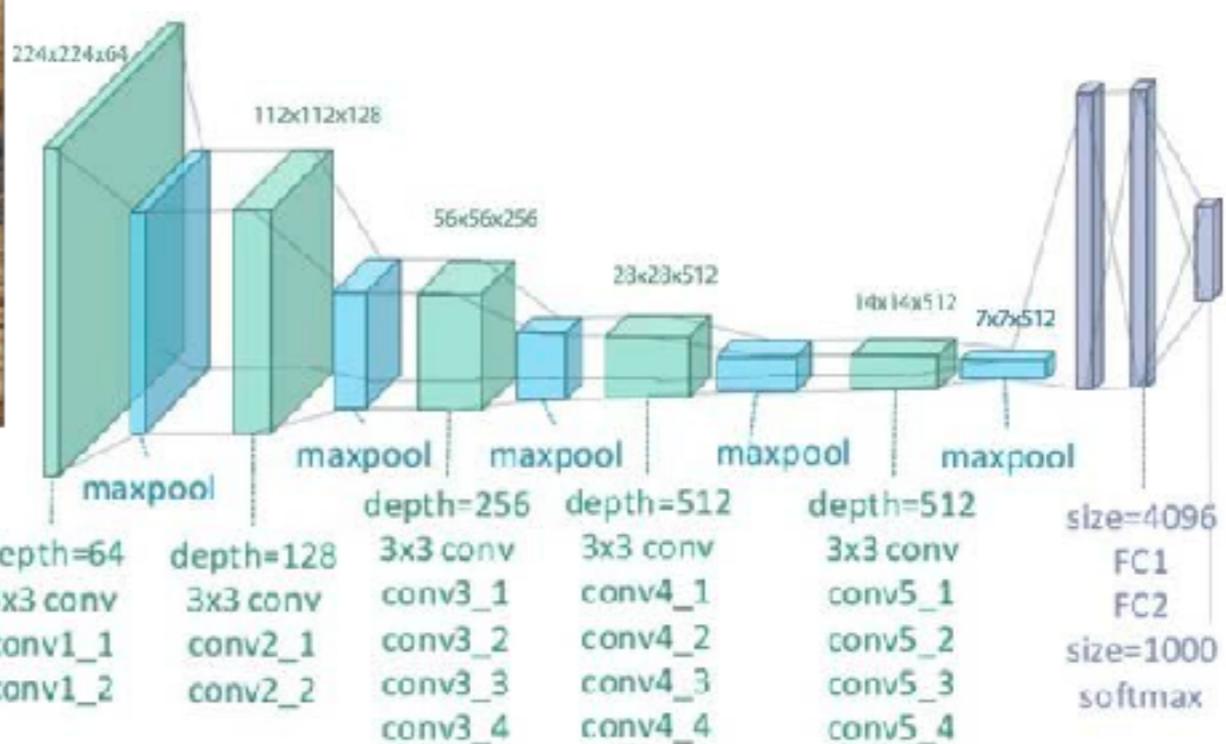
```



The VGG19 Deep Neural network  
(Image courtesy of Clifford K. Yang)

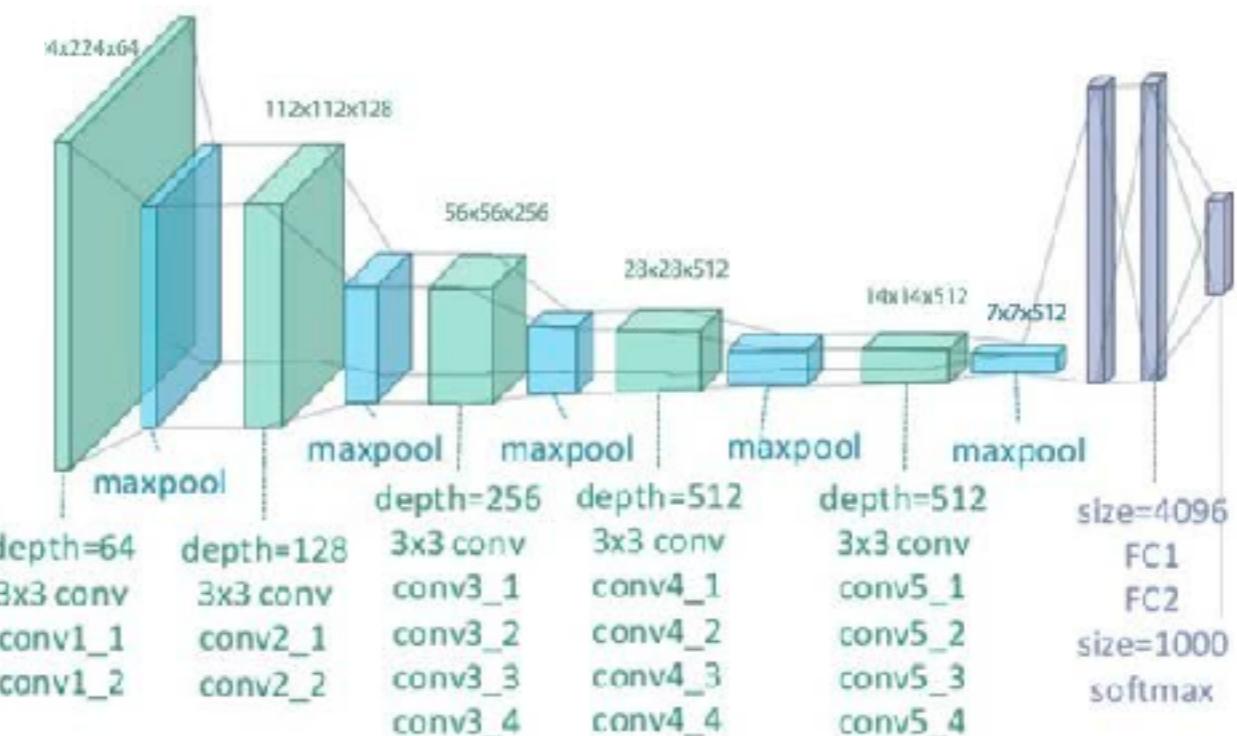
```
img = load("bicycle.jpg")
#show classify(vgg, img)
img

classify(vgg, img) = "mountain bike, all-terrain bike, off-roader"
```



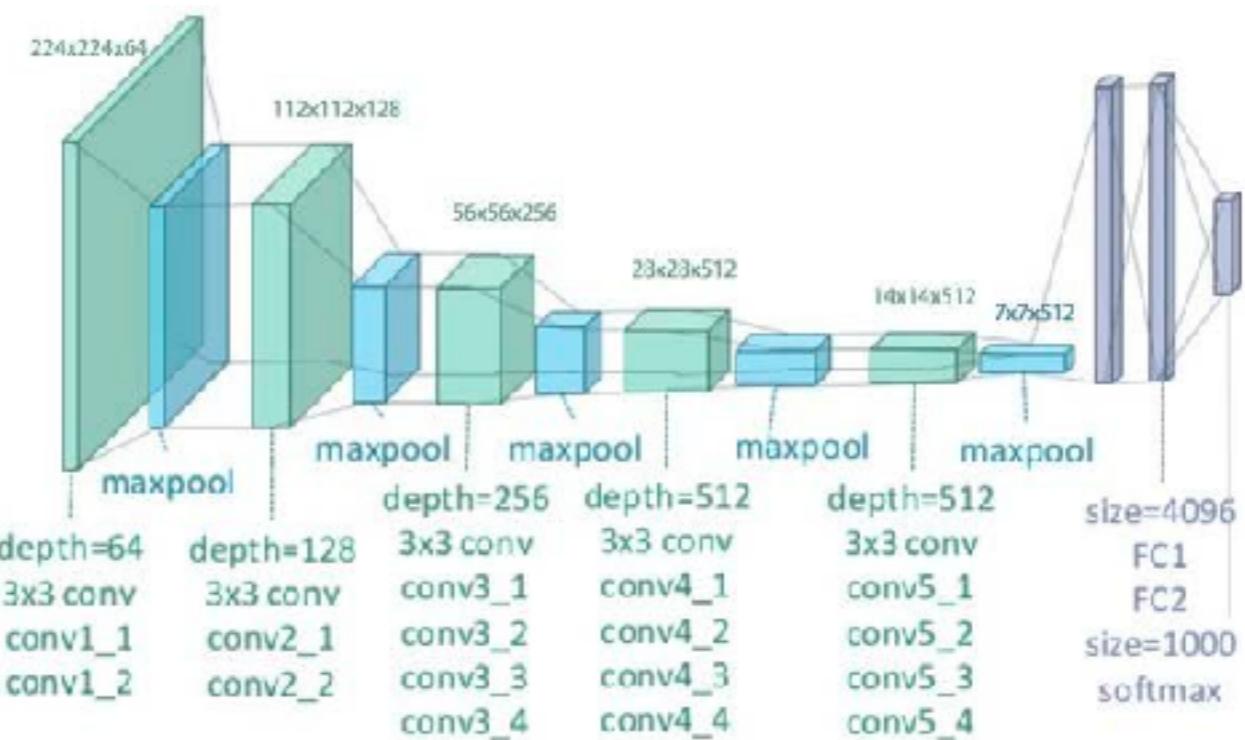
```
img = load("appleFruit.jpg")
@show classify(vgg,img)
img
```

```
classify(vgg, img) = "Granny Smith"
```



```
img = load("baby.jpg")
@show classify(vgg,img)
img
```

```
classify(vgg, img) = "diaper, nappy, napkin"
```



Data Science

Artificial Intelligence

Deep Learning

Machine Learning

Computer Science

Statistical Learning

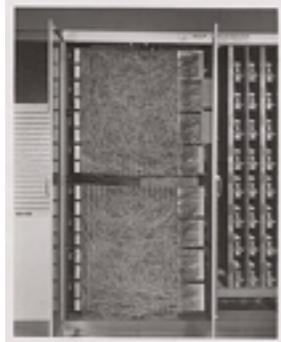
Data Mining

Statistics

Neuroscience



# Timeline of weak (narrow) AI



1950 Turing Test Proposed



1960's Neural Networks



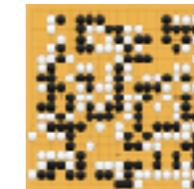
1970 Reverse Mode Automatic Differentiation



1980's Reinforcement Learning



~2000 Revival of machine learning

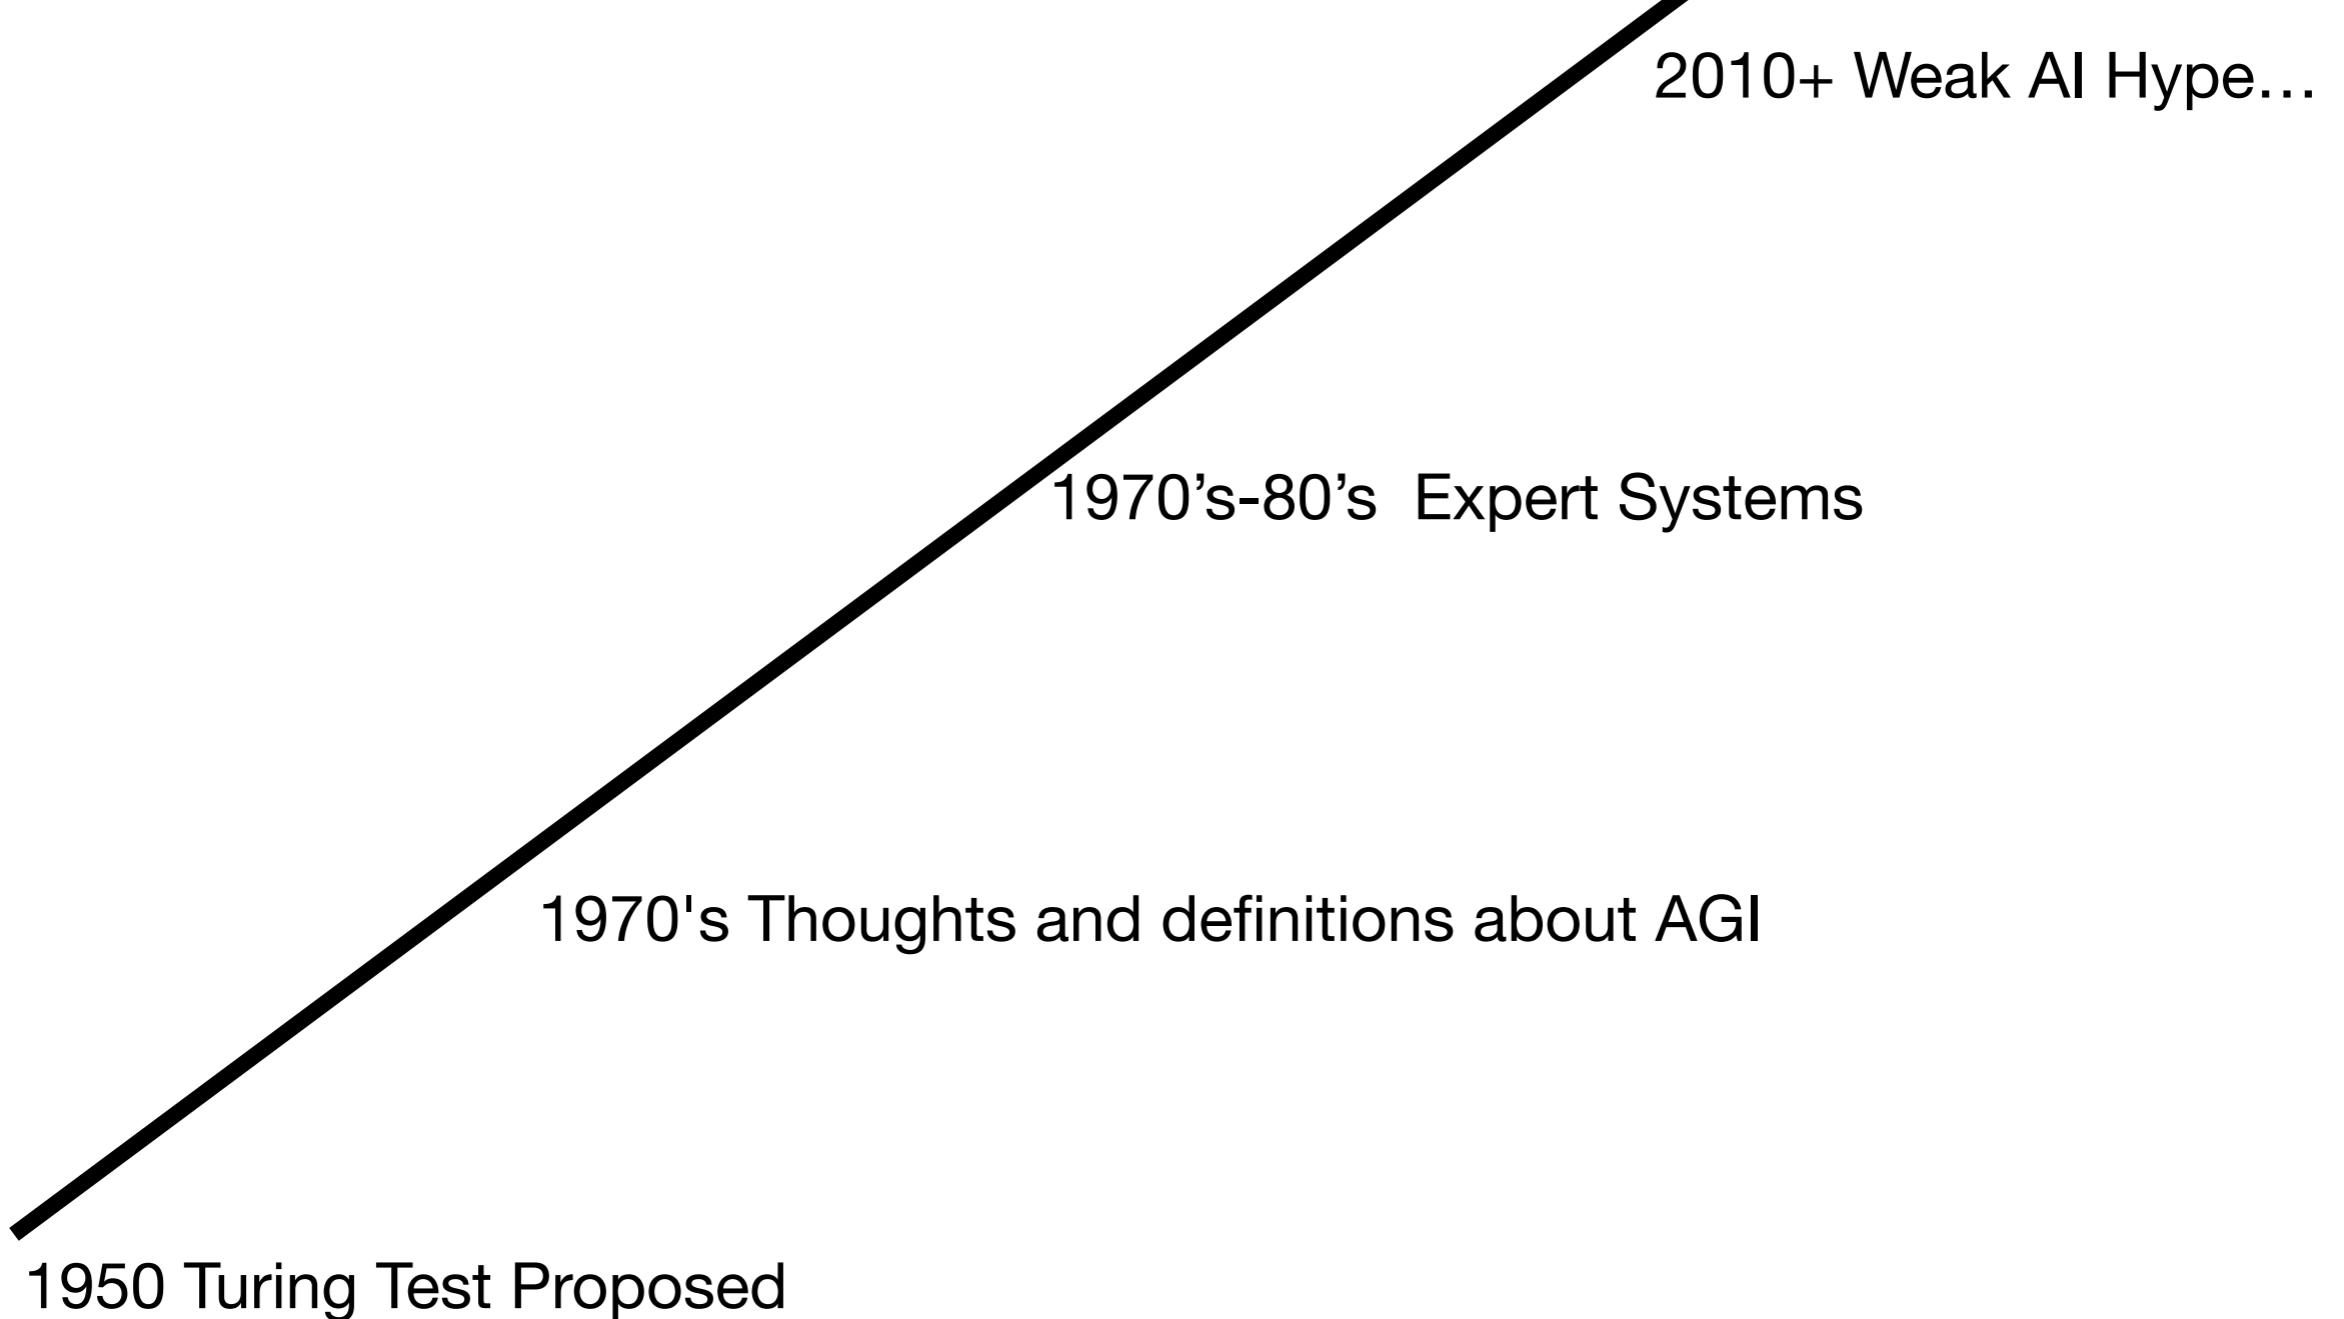


~2010 Machine Learning Renaissance begins

Weak AI  
‘Explosion’ and Deep Learning

# **Timeline of strong AI**

## **(Artificial General Intelligence)**



# Timeline of Deep Learning (prehistory)

“Deep Learning” is a thing!  
Deep Learning = AI?

2012 (**AlexNet**): ImageNet Classification with Deep Convolutional Neural Networks,  
by Alex Krizhevsky, Ilya Sutskever and Geoffrey E. Hinton.

1998: (**Convolutional nets**): Gradient based learning applied to document recognition,  
by Yann Lecun, Leon Bottou, Yoshua Bengio and Patrick Haffne.

1982 (**Hopfield nets**): Neural networks and physical systems with emergent collective computational abilities,  
by John Hopfield.

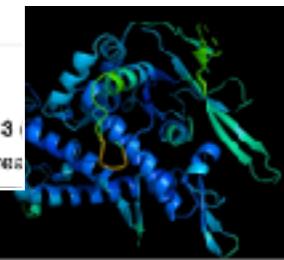
1958 (**Perceptron**): The perceptron: a probabilistic model for information storage and organization in the brain,  
by Frank Rosenblatt.

# Current Deep Learning

GPT-3

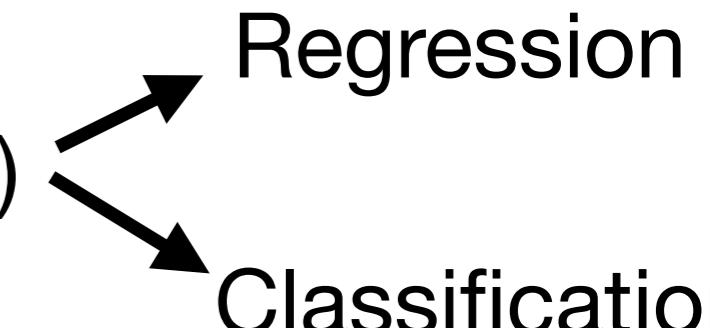
From Wikipedia, the free encyclopedia

Generative Pre-trained Transformer 3 is a large language model in the GPT series (and the successor to GPT-2) created by the company



- 2017 (**Transformers**): Attention is all you need, by Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan Gomez, Łukasz Kaiser and Illia Polosukhin.
- 2016 (**Deep RL vs. Go**): Mastering the game of Go with deep neural networks and tree search, by David Silver, Aja Huang, Chris Maddison and others.
- 2015 (**ResNet**): Deep residual learning for image recognition, by Kaiming He, Xiangyu Zhang, Shaoqing Ren and Jian Sun.
- 2015 (**Inception**): Going deeper with convolutions, by Christian Szegedy et. al.
- 2015 (**VGG**): Very deep convolutional networks for large-scale image recognition, by Karen Simonyan and Andrew Zisserman.
- 2014: (**GAN**): Generative adversarial nets, by Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville and Yoshua Bengio.
- 2014: (**ADAM**): Adam: A method for stochastic optimization, by Diederik Kingma and Jimmy Ba.
- 2013: (**Deep RL**): Playing Atari with Deep Reinforcement Learning by Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra and Martin Riedmiller.
- 2013: (**Inner layers**): Visualizing and Understanding Convolutional Networks, by Matthew Zeiler and Rob Fergus
- 2012 (**AlexNet**): ImageNet Classification with Deep Convolutional Neural Networks, by Alex Krizhevsky, Ilya Sutskever and Geoffrey E. Hinton.

# The main activities of machine learning

- Supervised learning ( $Y$  = labels,  $X$  = features)
  - Unsupervised learning (only  $X$ )
  - Semi-supervised learning (hybrid)
  - Reinforcement Learning (a time component)
  - Generative modeling (generating new  $X$ 's)
- 
- The diagram shows two arrows originating from the 'Supervised learning' item in the list. One arrow points to the word 'Regression' and the other points to the word 'Classification'. Both 'Regression' and 'Classification' are aligned vertically on the right side of the slide.

QQQQ: Note Yoni also on  
sequence data...

# **Part 2: The ML workflow and common “practice” data sets**

# The MNIST digits dataset

```
using Flux
imgs, labels = Flux.Data.MNIST.images(), Flux.Data.MNIST.labels();
@show length(imgs)
@show size(imgs[1])
@show labels[1:8]
imgs[1:8]
```

```
length(imgs) = 60000
size(imgs[1]) = (28, 28)
labels[1:8] = [5, 0, 4, 1, 9, 2, 1, 3]
```



# Basic EDA for MNIST

```
#Basic Exploratory Data Analysis (EDA) for MNIST
using Statistics, StatsPlots, Plots

x = heatmap(vcat(float.(im)...) for im in imgs)...)

d, n = size(x)
@show (d,n)

onMeanIntensity = mean(filter((u)->u>0,x))
@show onMeanIntensity

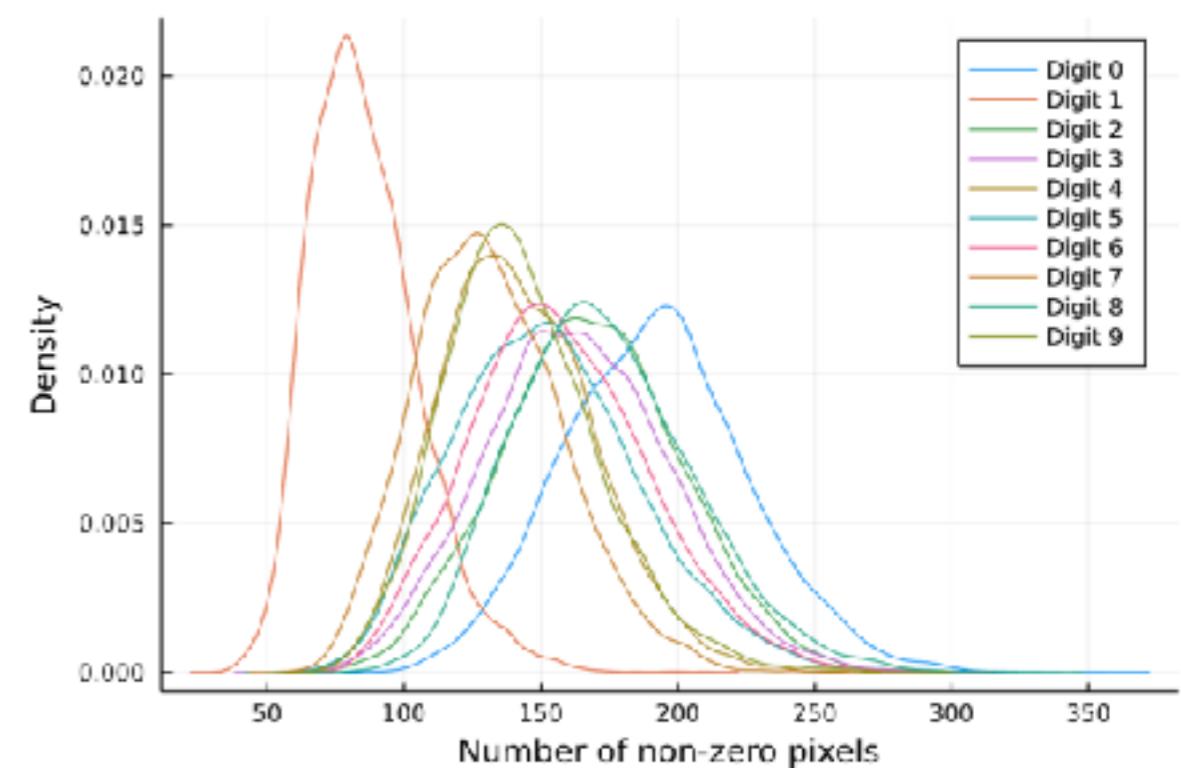
prop0 = sum(x .== 0)/(d*n)
@show prop0

prop1 = sum(x .== 1)/(d*n)
@show prop1

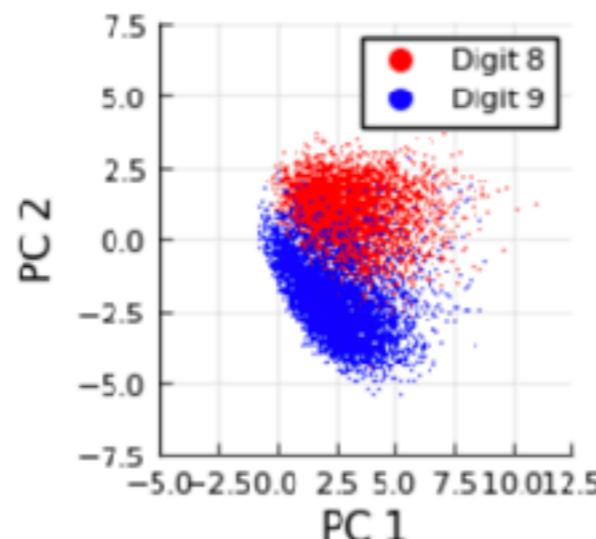
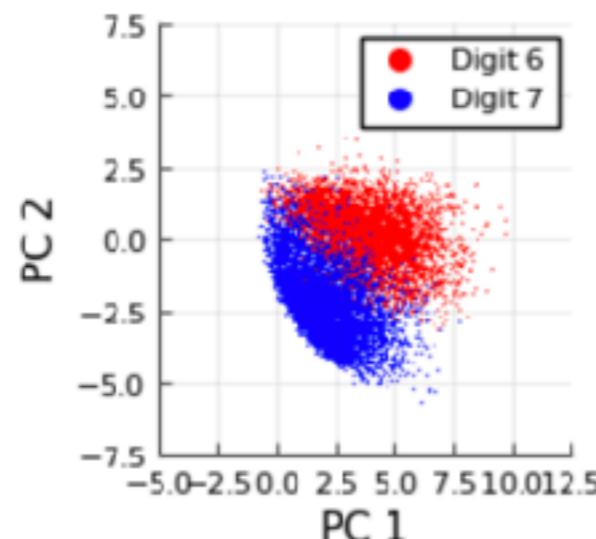
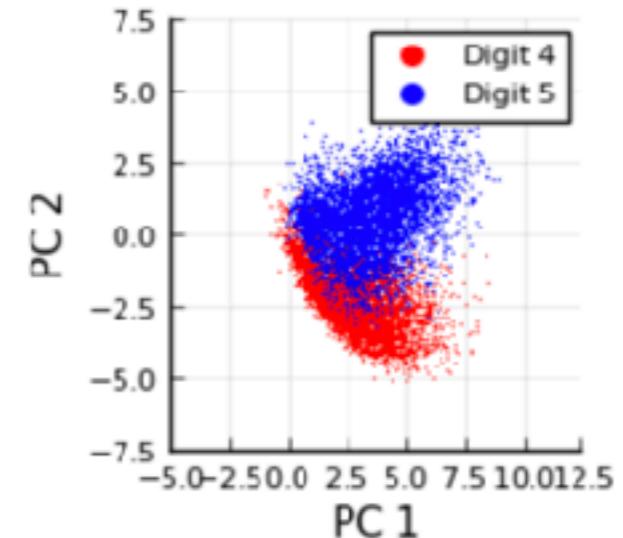
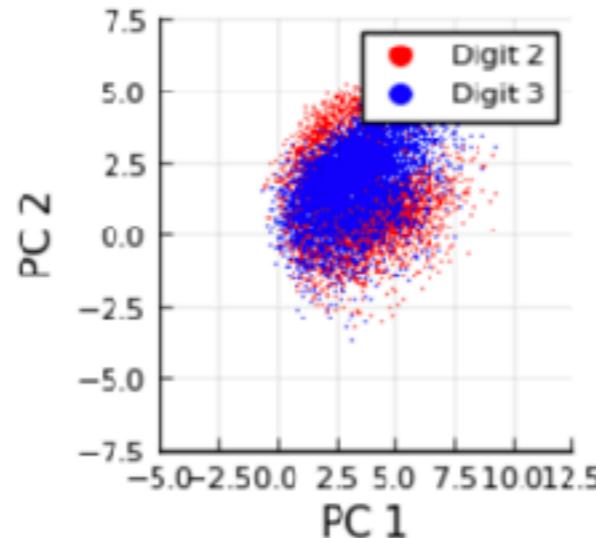
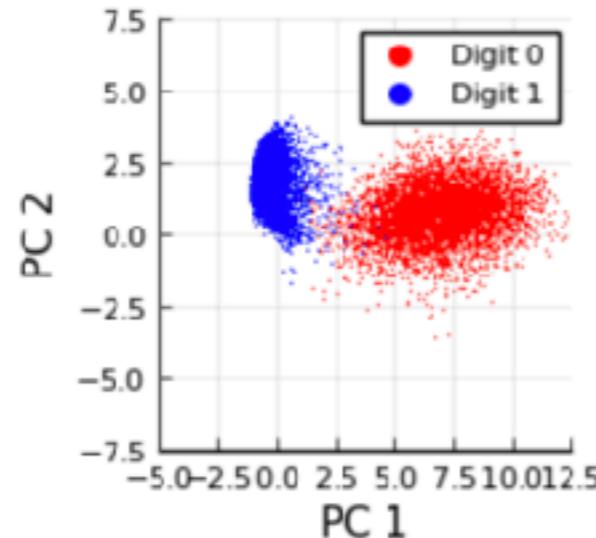
print("Label counts: ", [sum(labels .== k) for k in 0:9])

p = plot()
for k in 0:9
    onPixels = [ sum(x[:,i] .> 0) for i in (1:n)[labels .== k] ]
    p = density!(onPixels, label = "Digit $(k)")
end
plot(p,xlabel="Number of non-zero pixels", ylabel = "Density")

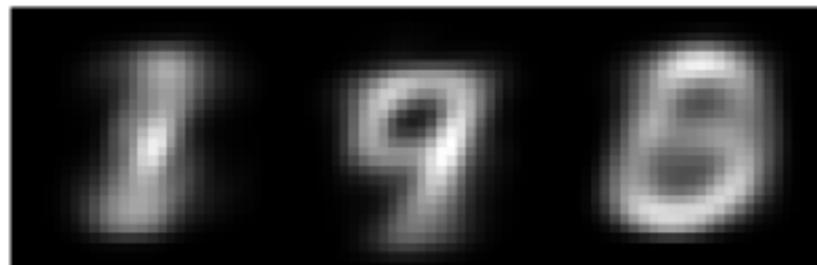
(d, n) = (784, 60000)
onMeanIntensity = Gray{Float32}(0.6933625f0)
prop0 = 0.8087977040816327
prop1 = 0.006681164965986395
Label counts: [5923, 6742, 5958, 6131, 5842, 5421, 5918, 6265, 5851, 5949]
```



# Breaking MNIST up via PCA (2 components)



# The centroids of k-means



k=3



k=5

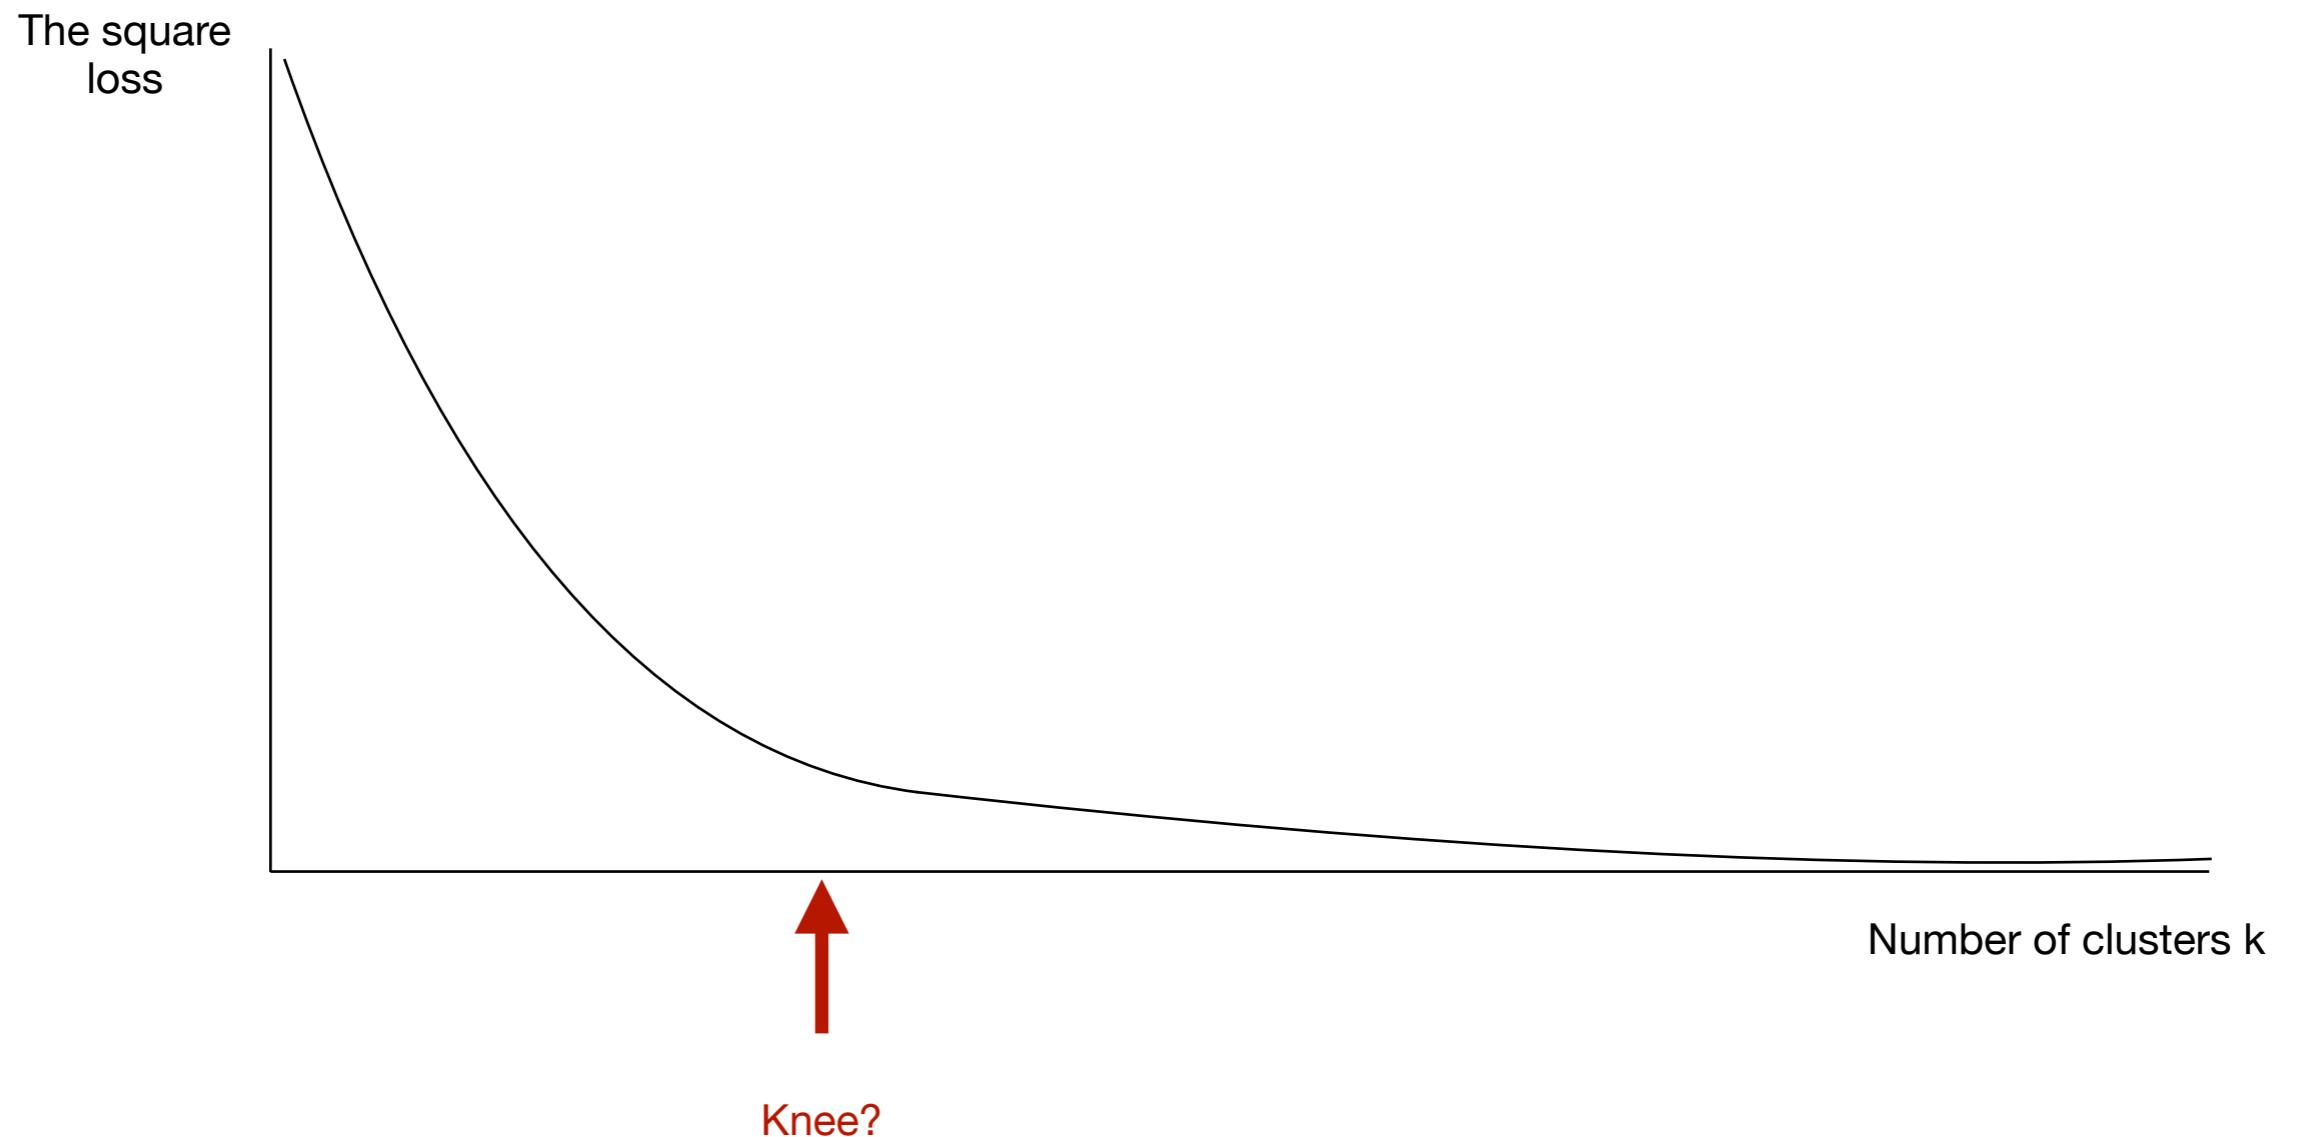


k=10



k=15

# Finding the “knee”



# A toy binary classifier digit 1, or not

```
: imgs[labels == 1][1:8]
```



```
: imgs[labels != 1][1:8]
```



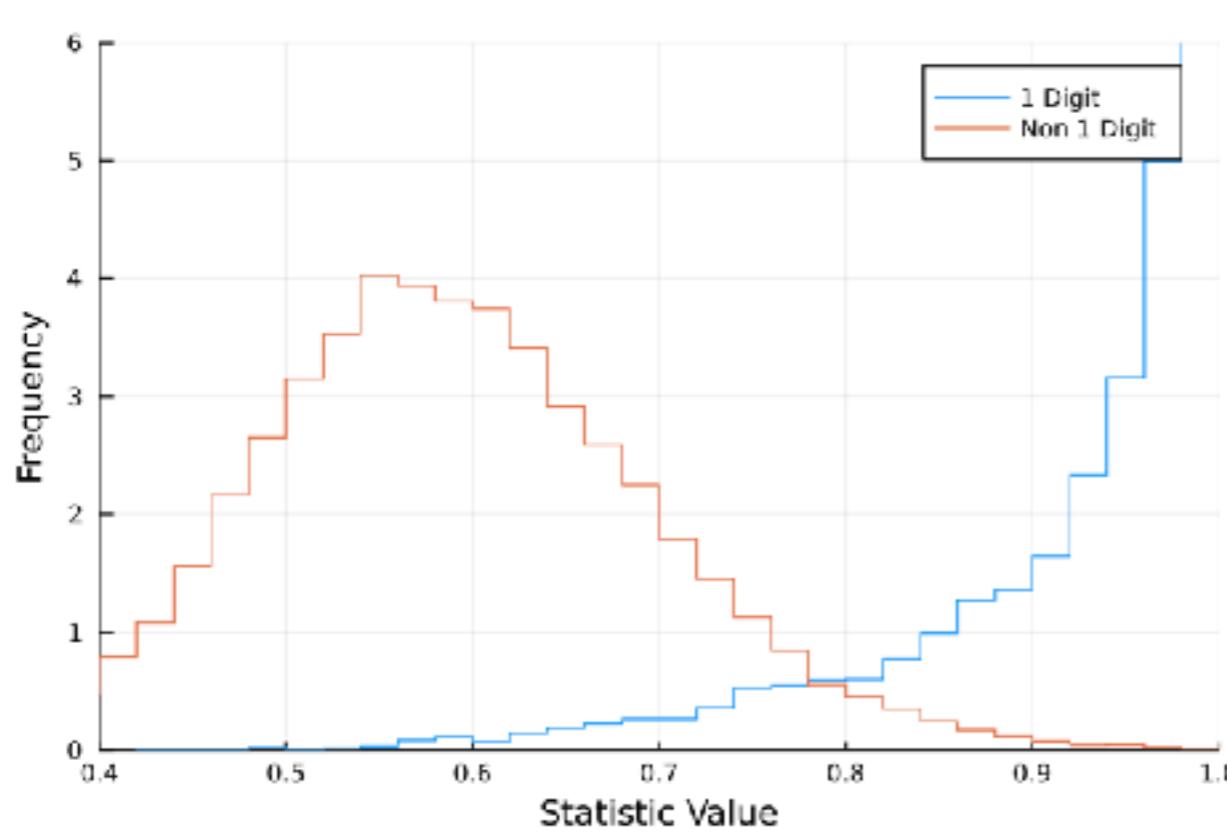
Define a statistic (specific to the digit 1) : Count the proportion of intensity around the peak per row:

$$\chi(x) = \frac{\sum_{i=1}^{28} \sum_{k=-2}^2 x_{i,m(x,i)+k}}{\sum_{i=1}^{28} \sum_{j=1}^{28} x_{i,j}}, \quad \text{where } m(x, i) = \operatorname{argmax}_{j=1, \dots, 28} x_{i,j}$$

A classifier:  $\hat{f}_\theta(x) = \begin{cases} -1 & \chi(x) \leq \theta, \\ +1 & \chi(x) > \theta. \end{cases}$

# A toy binary classifier digit 1, or not

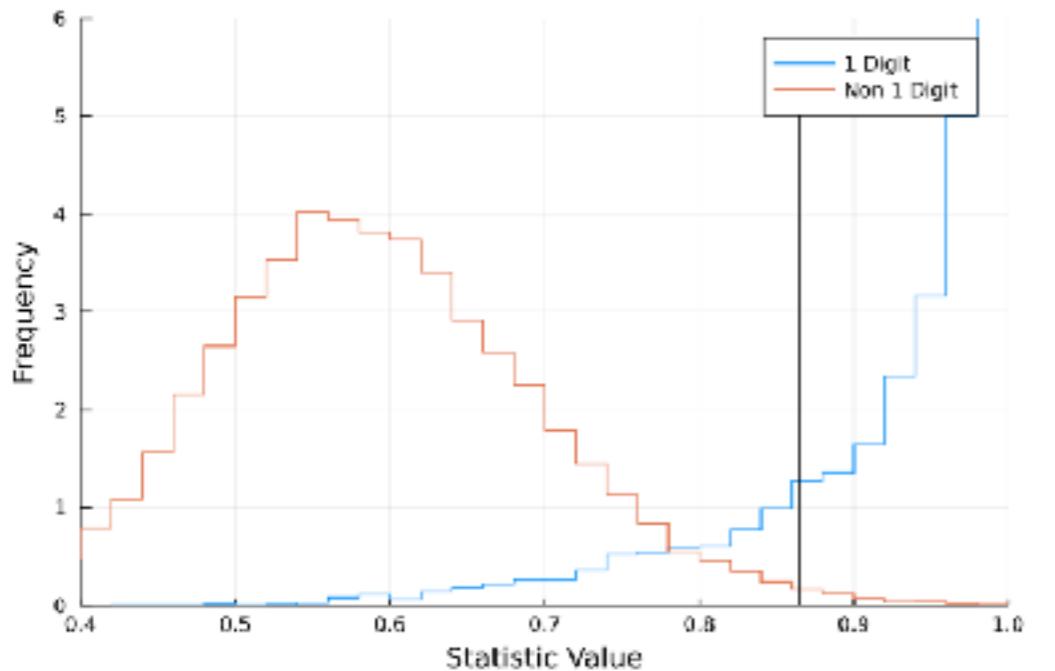
$$\chi(x) = \frac{\sum_{i=1}^{28} \sum_{k=-2}^2 x_{i,m(x,i)+k}}{\sum_{i=1}^{28} \sum_{j=1}^{28} x_{i,j}}, \quad \text{where} \quad m(x, i) = \operatorname{argmax}_{j=1, \dots, 28} x_{i,j}$$



$$\hat{f}_\theta(x) = \begin{cases} -1 & \chi(x) \leq \theta, \\ +1 & \chi(x) > \theta. \end{cases}$$

How to choose the threshold  $\theta$ ?

# Say we chose it at $\theta = 0.865$



		Decision	
		Decide $-1$ (8,916)	Decide $+1$ (1,084)
Reality	Label is $-1$ (8,865)	True negative (8,781)	False positive (84)
	Label is $+1$ (1,035)	False negative (135)	True Positive (1000)

# Say we chose it at $\theta = 0.865$

		Decision	
		Decide $-1$ (8,916)	Decide $+1$ (1,084)
Reality	Label is $-1$ (8,865)	True negative (8,781)	False positive (84)
	Label is $+1$ (1,035)	False negative (135)	True Positive (1000)

$$\text{Precision} = \frac{|\text{true positive}|}{|\text{true positive}| + |\text{false positive}|},$$

$$\text{Recall} = \frac{|\text{true positive}|}{|\text{true positive}| + |\text{false negative}|}.$$

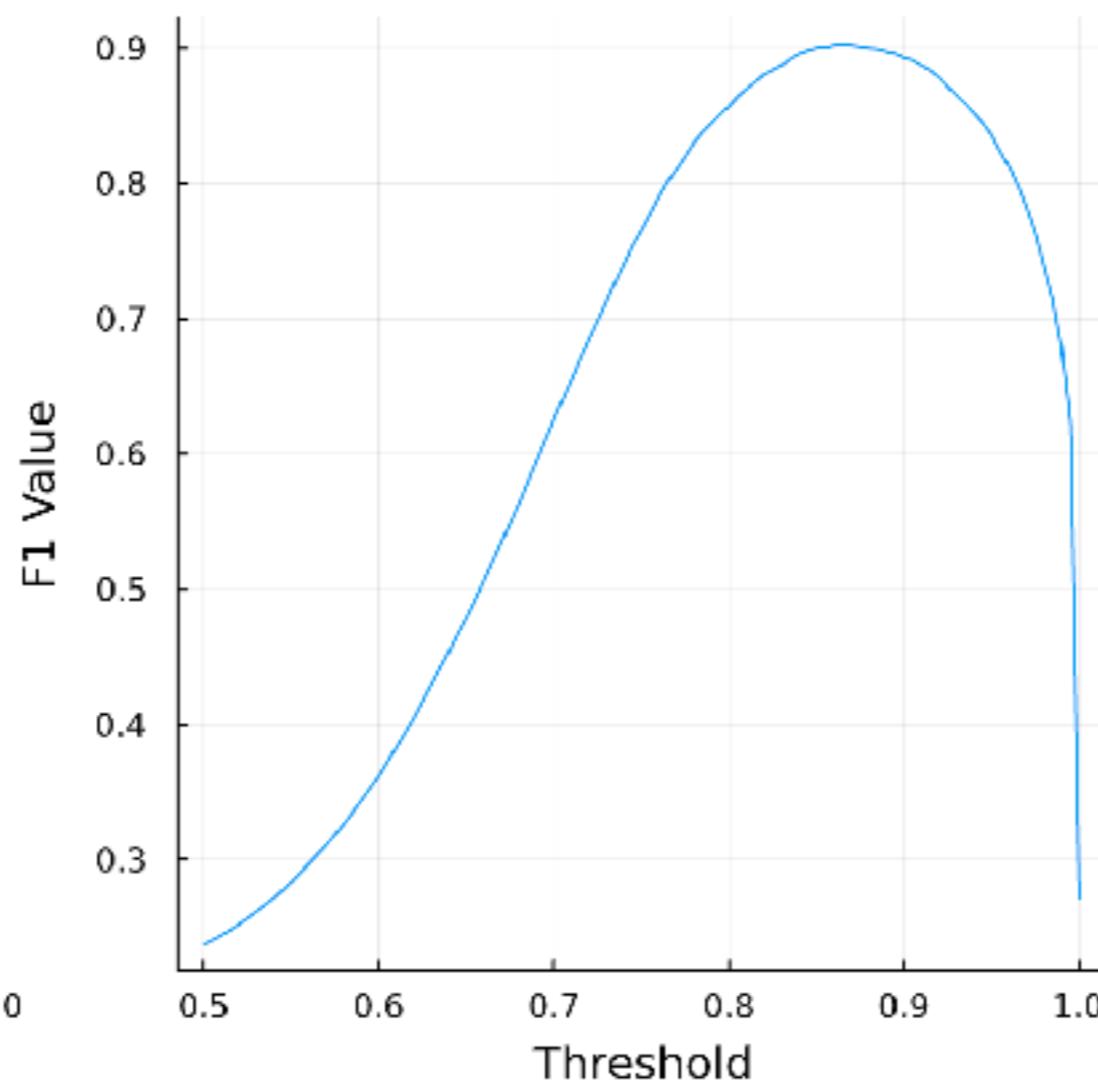
$$\text{Precision} = \frac{1000}{1000 + 84} = 92.25\%,$$

$$\text{Recall} = \frac{1000}{1000 + 135} = 88.11\%.$$

$$F_1 = \frac{2}{\frac{1}{\text{Precision}} + \frac{1}{\text{Recall}}} = 2 \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} = 90.13\%$$

# How we chose $\theta = 0.865$

$$F_1 = \frac{2}{\frac{1}{\text{Precision}} + \frac{1}{\text{Recall}}} = 2 \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} = 90.13\%$$



# Why not just accuracy?

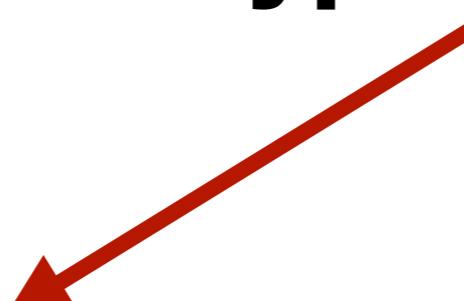
$$\text{accuracy} = \frac{1}{m} \sum_{i=1}^m \mathbf{1} \left\{ \hat{f}(x^{(i)}) = y_i \right\}$$

## What can be a problem?

# Recap:

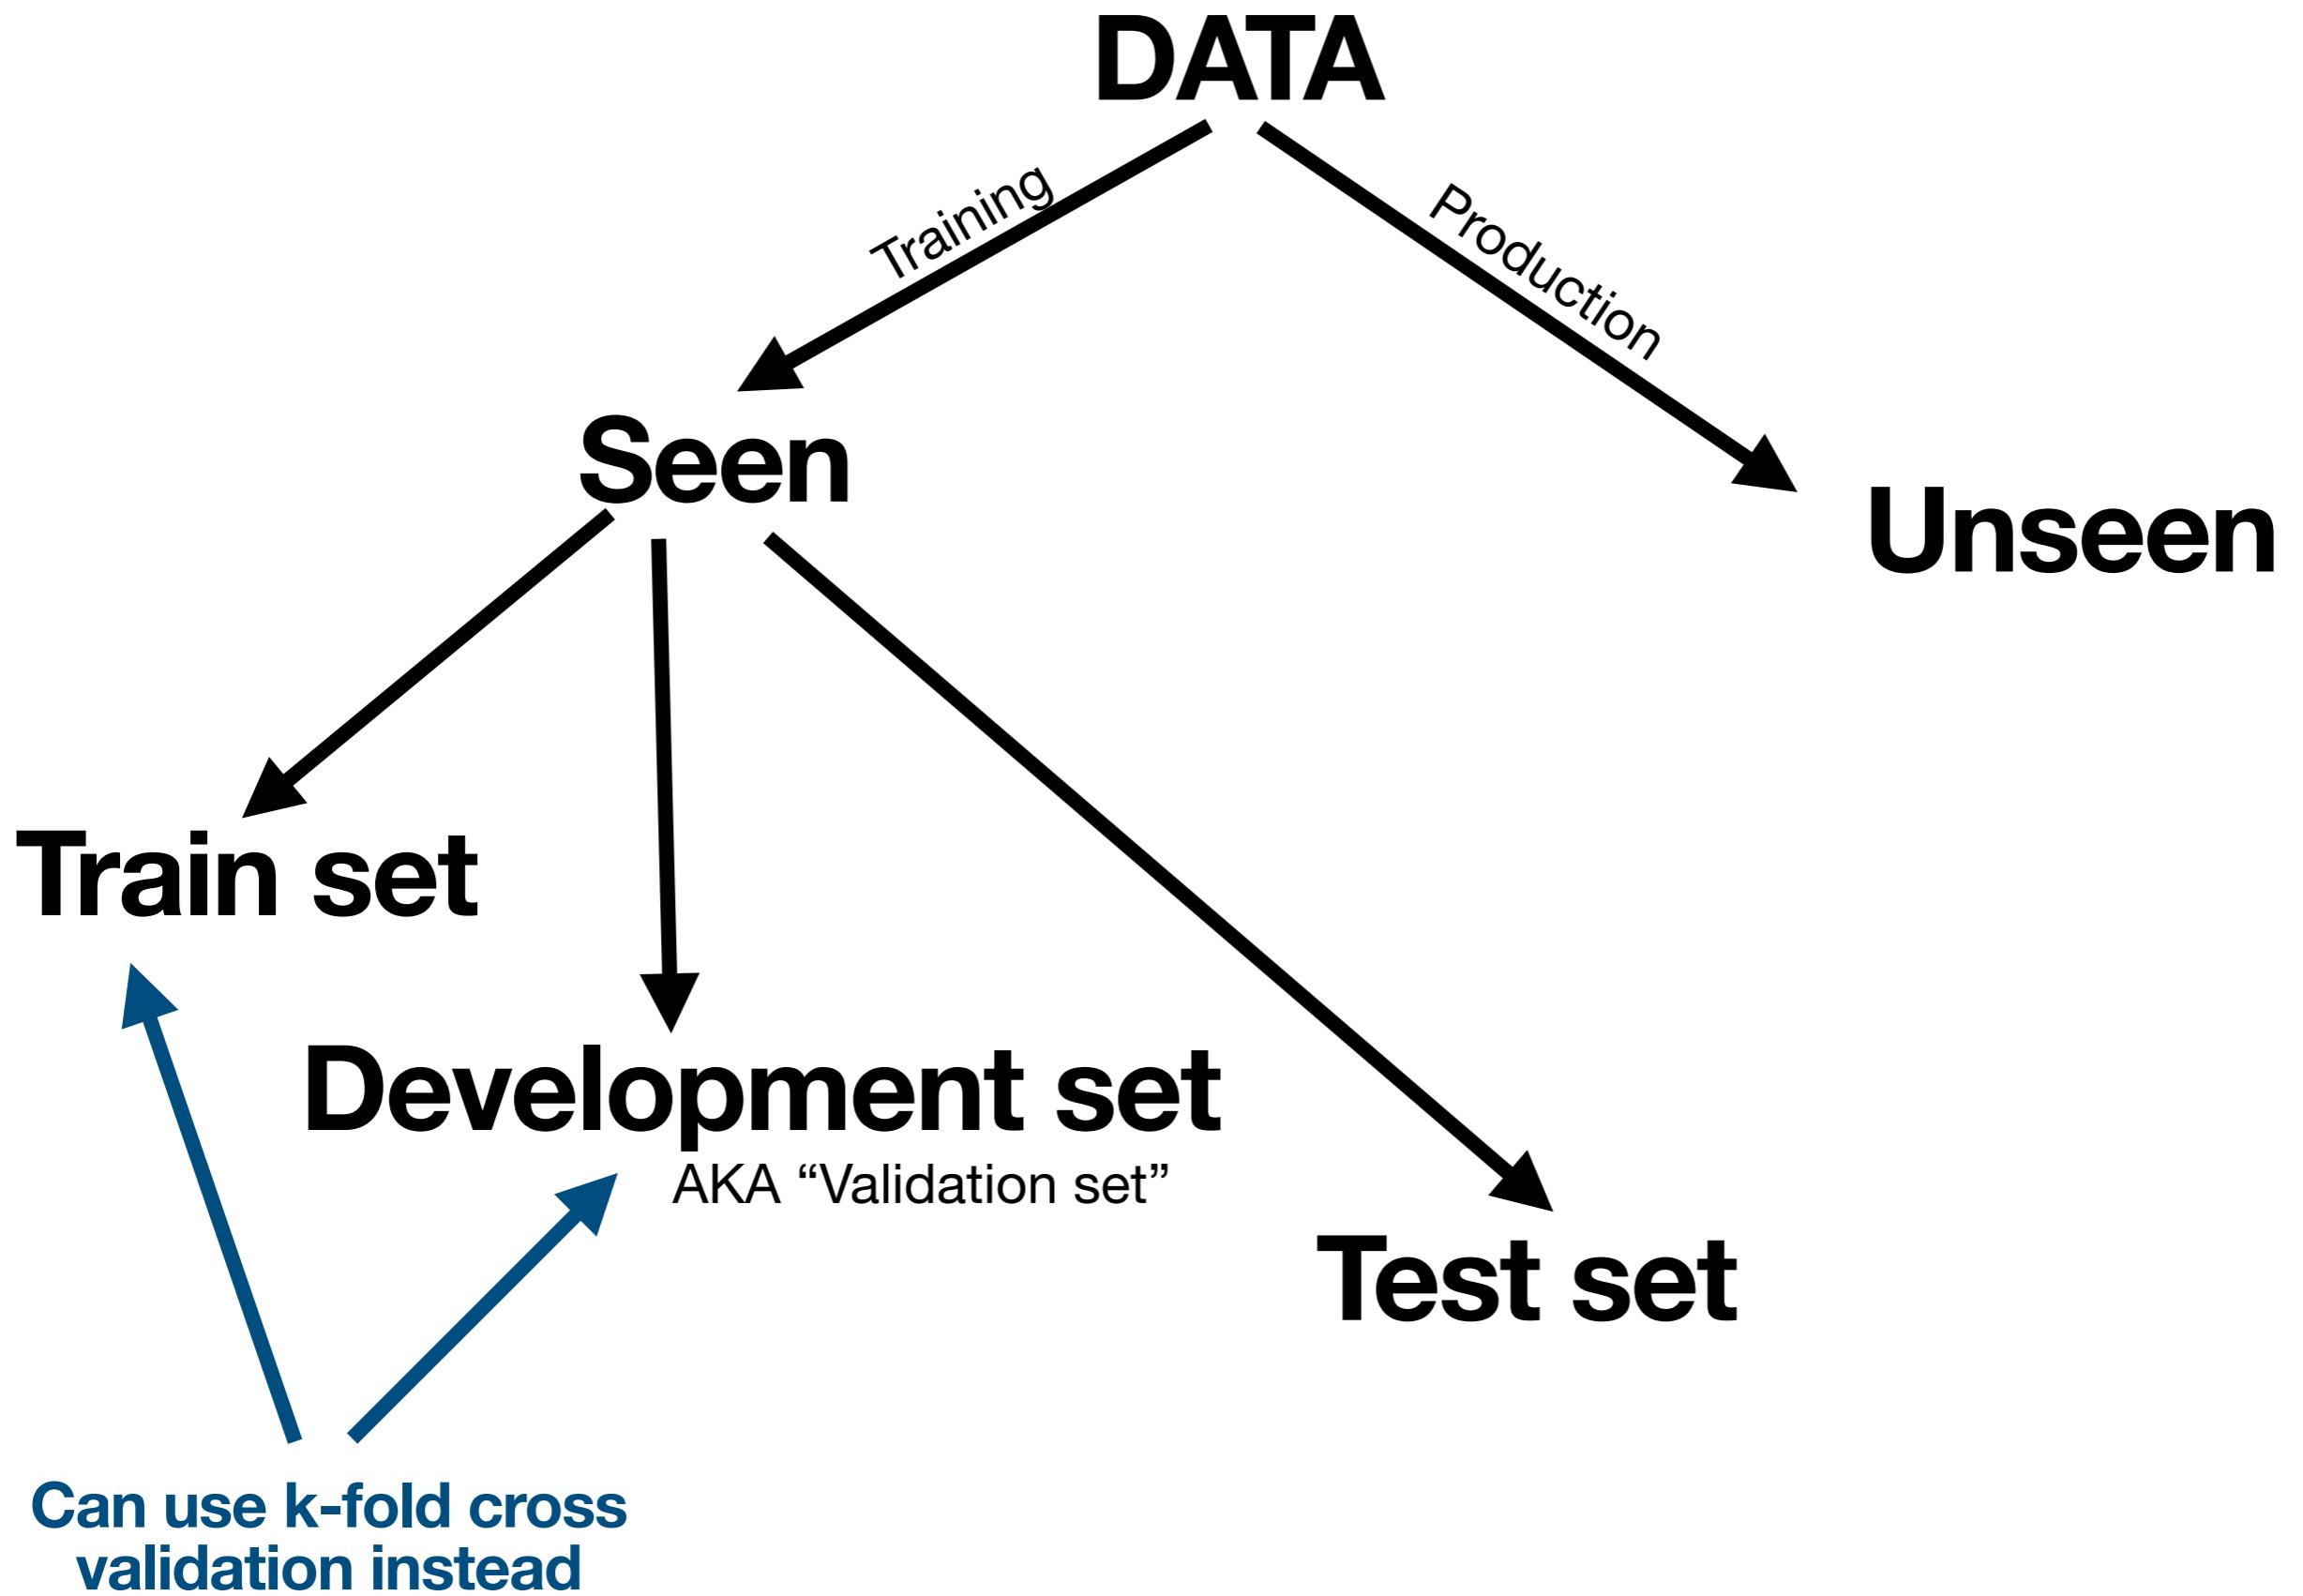
We “learned” the parameter  $\theta$

We had some hyper-parameters too:

$$\chi(x) = \frac{\sum_{i=1}^{28} \sum_{k=-2}^2 x_{i,m(x,i)+k}}{\sum_{i=1}^{28} \sum_{j=1}^{28} x_{i,j}}, \quad \text{where } m(x, i) = \operatorname{argmax}_{j=1,\dots,28} x_{i,j}$$


How do we choose hyper-parameters?

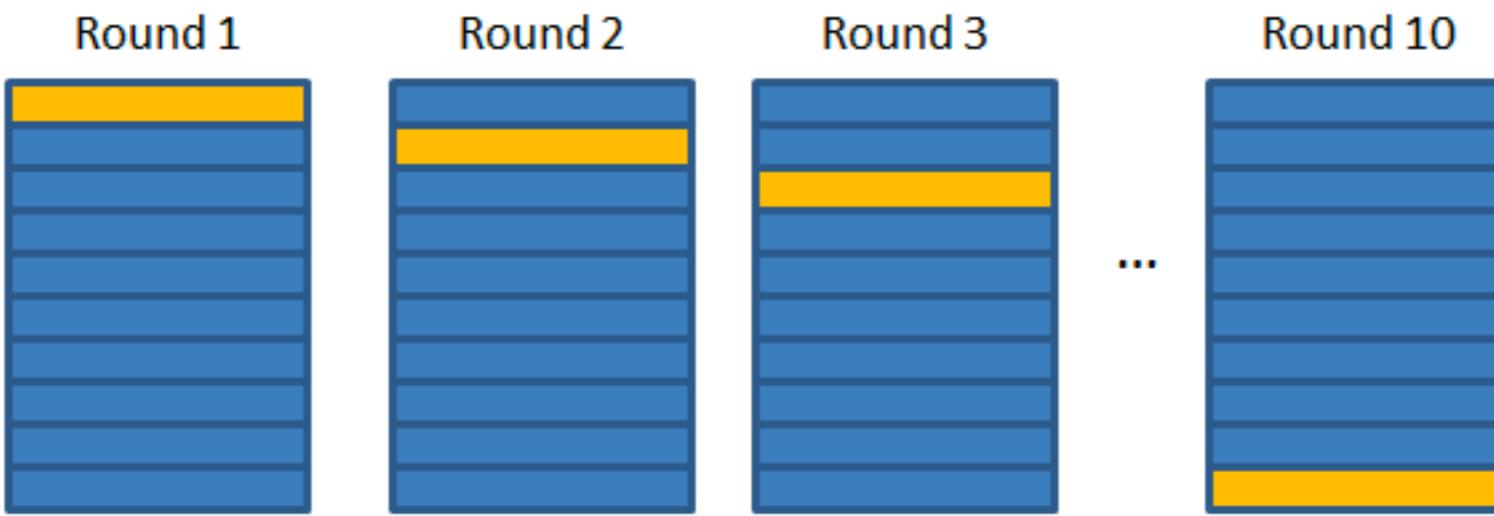
How do we test our classifier?



# K-fold cross validation

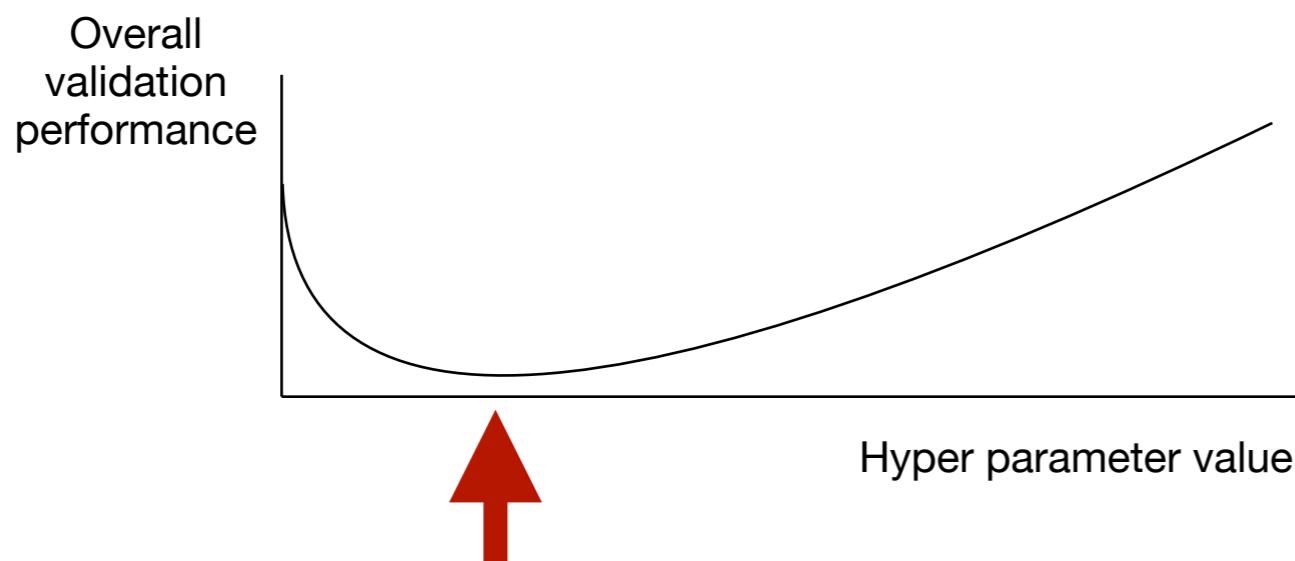
Validation Set  
Training Set

k=10



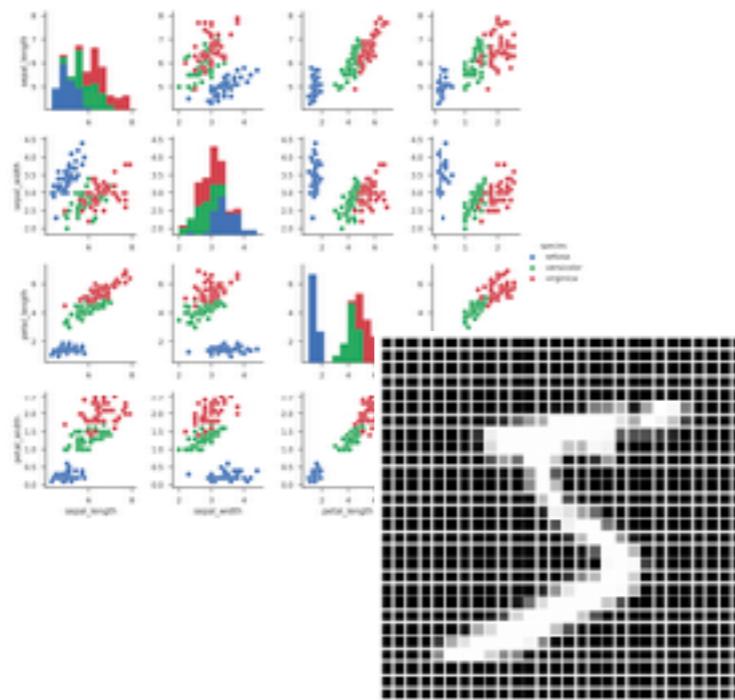
Validation performance:      87.3%      92.5%      91.2%      89.6%

Overall validation performance:  $\text{mean}(87.3, 92.5, 91.2, \dots, 89.6)$

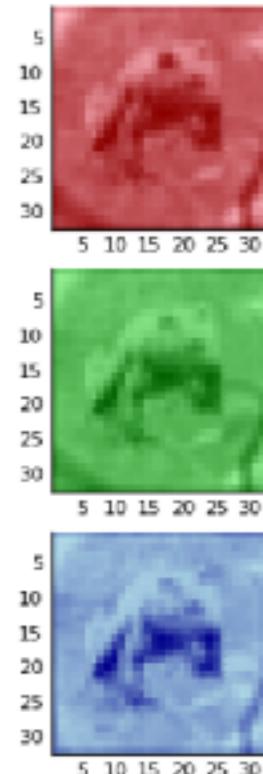
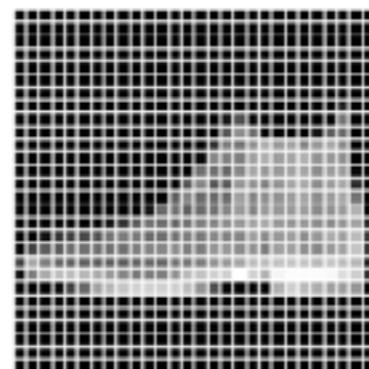


# Some more popular “toy” datasets

1.Iris Dataset



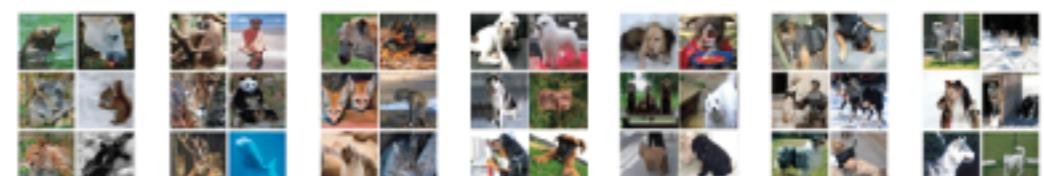
2.MNIST



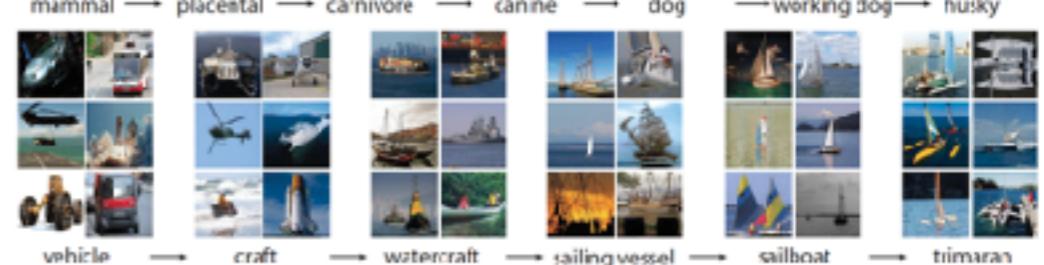
3.Fashion MNIST



4.CIFAR-10



5.ImageNet



6.Twitter Sentiment Analysis



# **Part 3: Some tools you know - used for ML**

# A linear (binary) classifier

Positive (+1)

```
imgs[labels == 3][1:8]
```

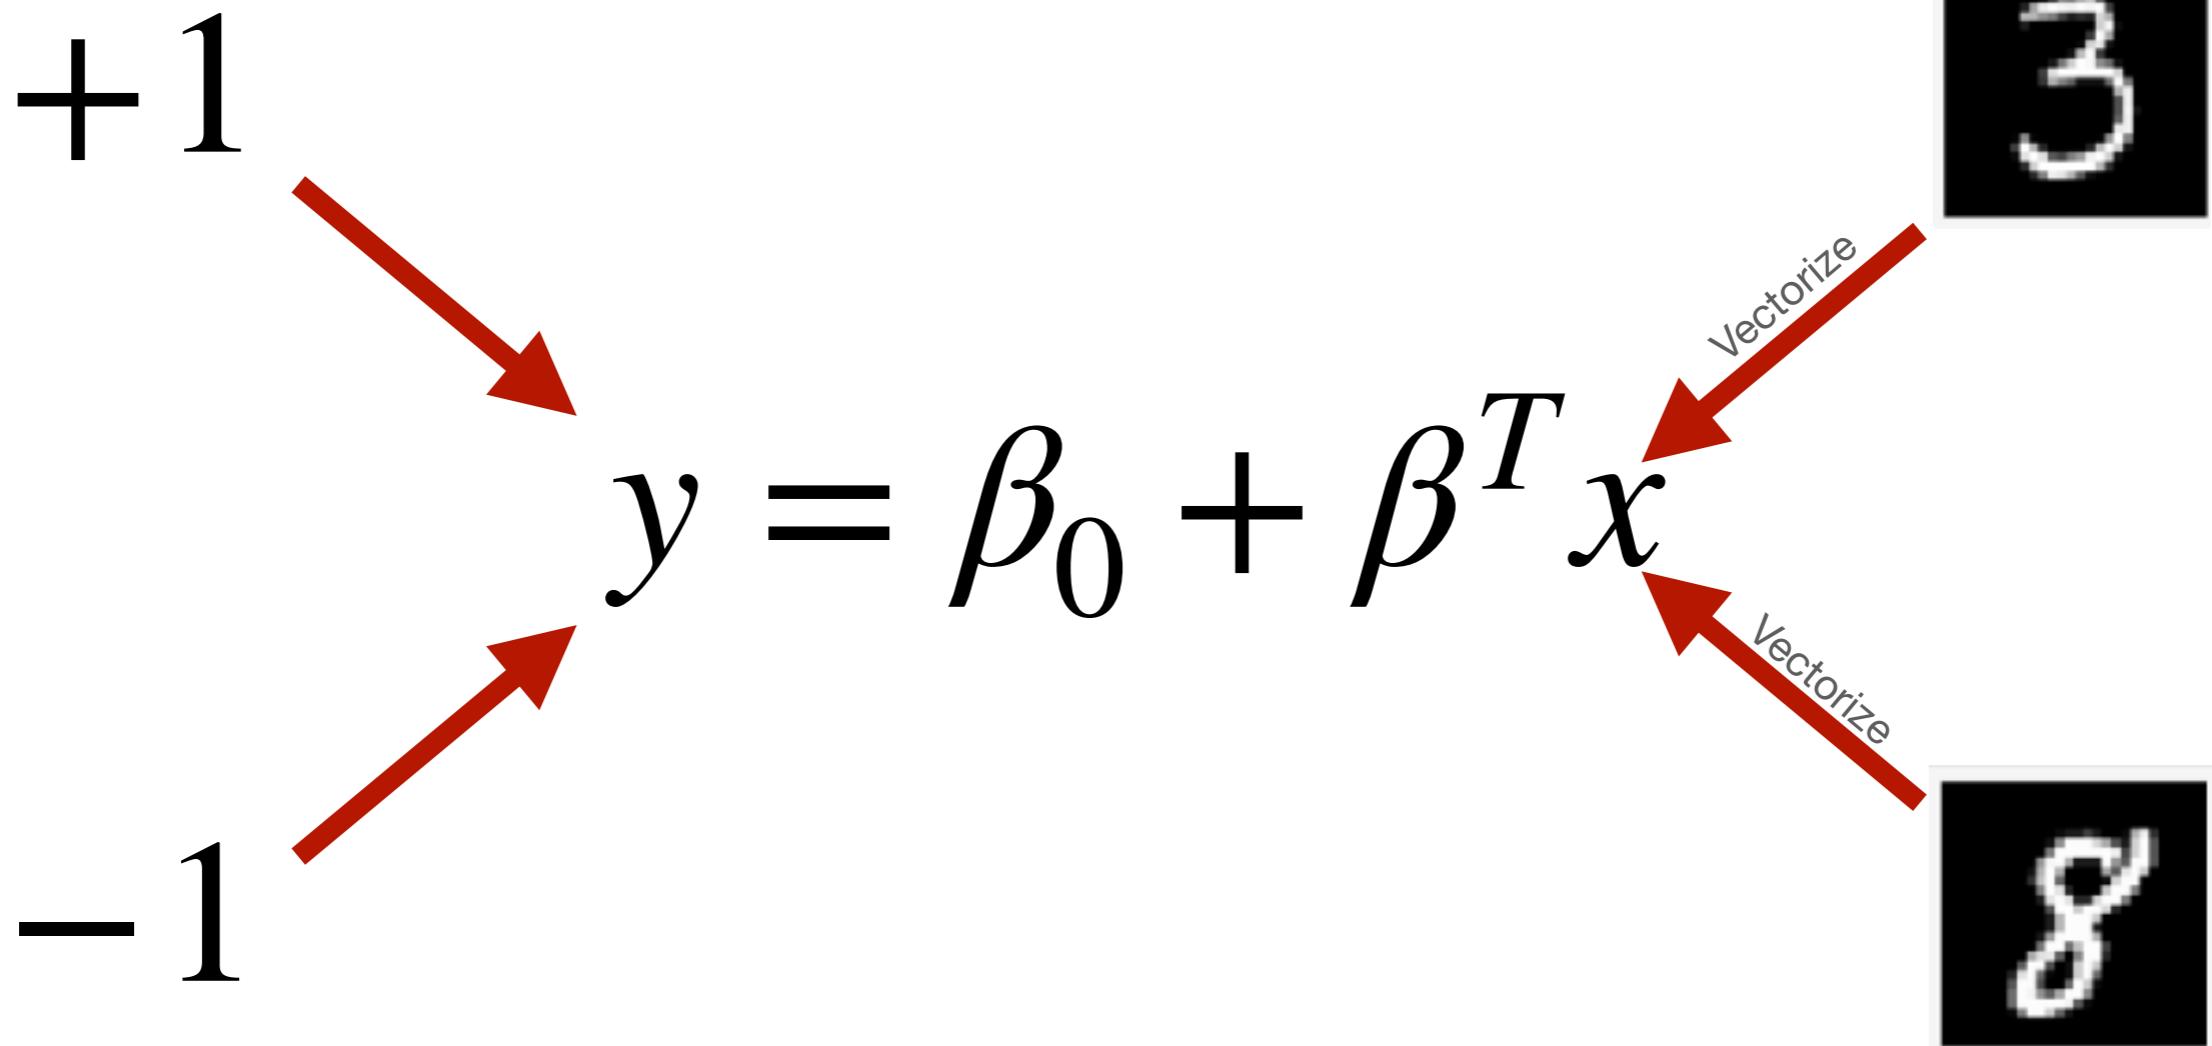


Negative (-1)

```
imgs[labels == 8][1:8]
```



# A linear (binary) classifier



Minimize:  $\text{Loss}(x, y) = \sum_{i=1}^m (y_i - \beta_0 - \beta^T x_i)^2$

Classifier:  $\hat{f}(x) = \text{sign}(\hat{\beta}_0 + \hat{\beta}^T x)$

# A linear (binary) classifier

Minimize:  $\text{Loss}(x, y) = \sum_{i=1}^m (y_i - \beta_0 - \beta^T x_i)^2 = \|y - A\beta\|^2$

Option 1:  $\hat{\beta} = A^\dagger y$       Sometimes  $\downarrow$   $A^\dagger = (A^T A)^{-1} A^T$

Option 2:  $\hat{\beta}(0), \hat{\beta}(1), \hat{\beta}(2), \hat{\beta}(3), \dots$       With (some form of) gradient descent

$$\hat{\beta}(t+1) = \hat{\beta}(t) - \eta \nabla L(\hat{\beta}(t))$$

$$\nabla L(\beta) = 2A^T(A\beta - y)$$

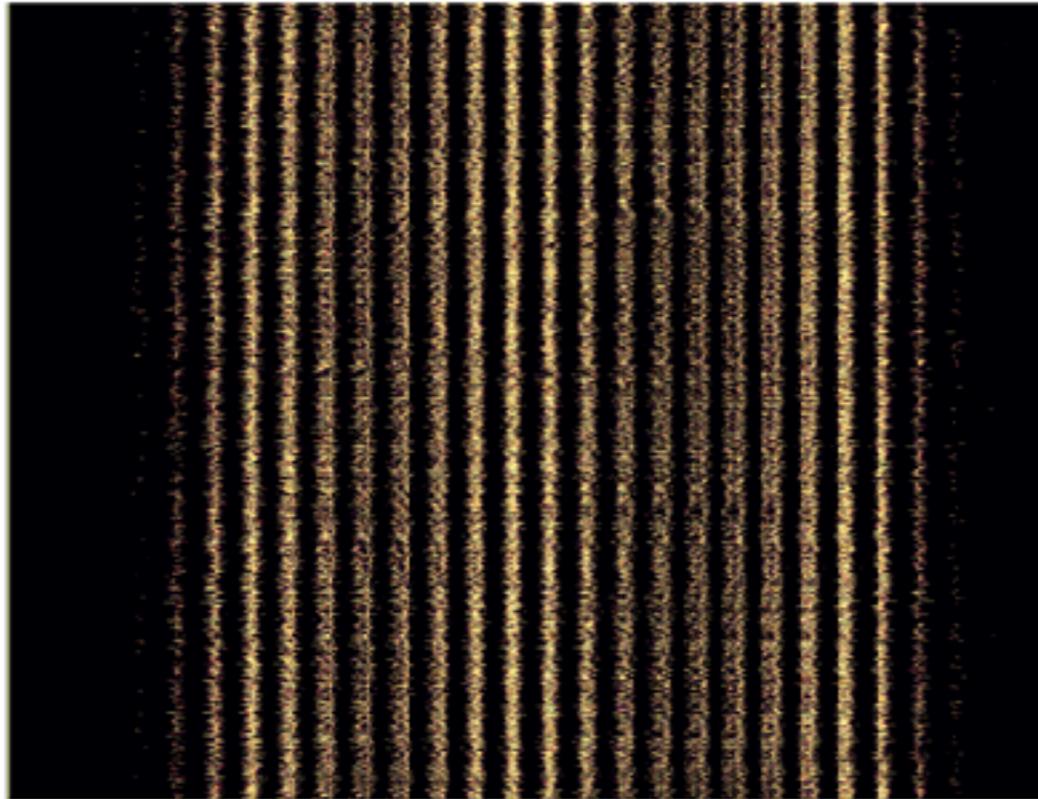
# Let's do it!

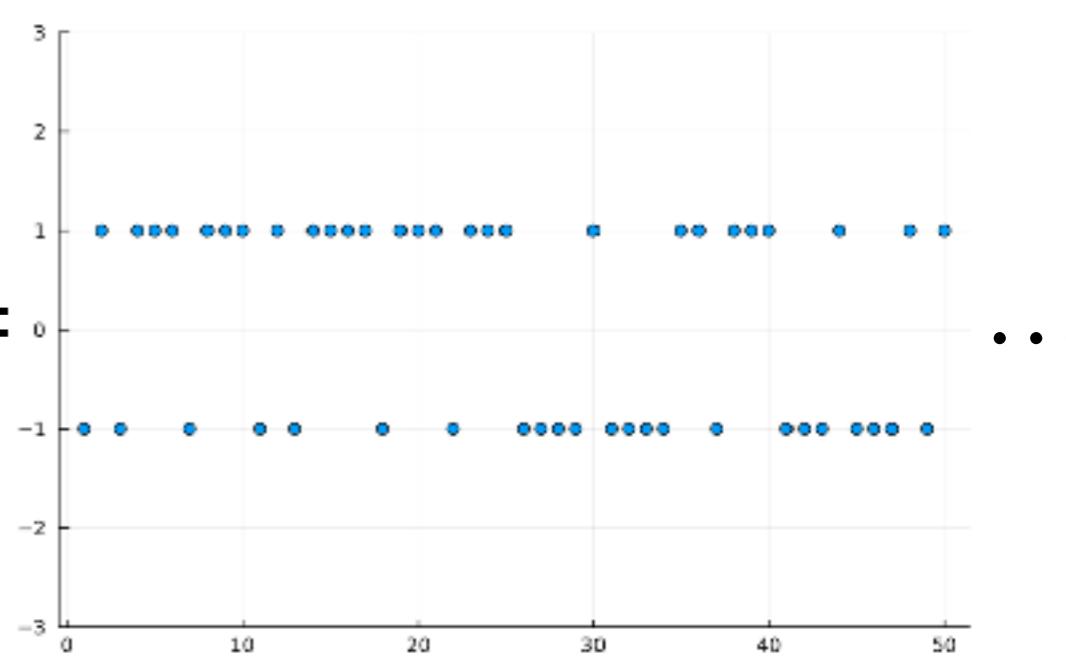
Option 1:  $\hat{\beta} = A^\dagger y$

Sometimes

$$A^\dagger = (A^T A)^{-1} A^T$$

Option 2:  $\hat{\beta}(t + 1) = \hat{\beta}(t) - \eta 2A^T(A\hat{\beta}(t) - y)$

$$A =$$


$$y =$$


# **Activity A: MNIST digit classification with least squares**

# Exercise: Work on these datasets

Task 1a:  $\hat{\beta} = A^\dagger y$

Task 1b:  $\hat{\beta}(t + 1) = \hat{\beta}(t) - \eta 2A^T(A\hat{\beta}(t) - y)$

Task 2: Evaluate test accuracy

<https://github.com/yoninazarathy/MathematicalEngineeringDeepLearning/blob/master/ML180Minutes/exercise.r>

```
## ML in 180 Minutes group/self exercise...

#Download training design matrix and labels
A <- read.csv(url("https://raw.githubusercontent.com/yoninazarathy/MathematicalEngineeringDeepLearning/master/ML180Minutes/A_3vs8_matrix_train.csv"))
y <- read.csv(url("https://raw.githubusercontent.com/yoninazarathy/MathematicalEngineeringDeepLearning/master/ML180Minutes/y_3vs8_vector_train.csv"))

#Download testing design matrix (each row is an image) for the digit "3" (positive)
A3test <- read.csv(url("https://raw.githubusercontent.com/yoninazarathy/MathematicalEngineeringDeepLearning/master/ML180Minutes/A_3_test.csv"))

#Download testing design matrix (each row is an image) for the digit "8" (negative)
A8test <- read.csv(url("https://raw.githubusercontent.com/yoninazarathy/MathematicalEngineeringDeepLearning/master/ML180Minutes/A_8_test.csv"))

#Task 1: Finding the least squares beta based on A and y.
#Task 1a: Do this via the pseudoinverse, or GLM, or similar.
#Task 1b: Do this via an iterative gradient descent implementation
#Task 2: Evaluate the accuracy of the estimator(s) from task 1 on the accuracy|
```

# Multi-class: One vs. rest

```
using Flux.Data.MNIST, PyPlot, LinearAlgebra
using Flux: onehotbatch

imgs = MNIST.images()
labels = MNIST.labels()
nTrain = length(imgs)

trainData = vcat([hcat(float.(imgs[i])...) for i in 1:nTrain]...)
trainLabels = labels[1:nTrain];

testImgs = MNIST.images(:test)
testLabels = MNIST.labels(:test)
nTest = length(testImgs)

testData = vcat([hcat(float.(testImgs[i])...) for i in 1:nTest]...)

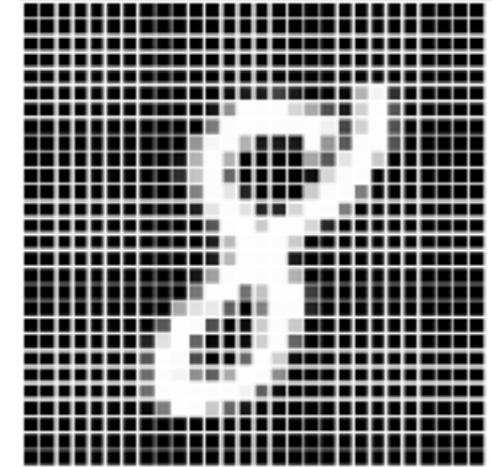
A = [ones(nTrain) trainData];
Adag = pinv(A);
tfPM(x) = x ? +1 : -1
yDat(k) = tfPM.(onehotbatch(trainLabels,0:9)'[:,k+1])
bets = [Adag*yDat(k) for k in 0:9];

classify(input) = findmax(([1 ; input]'*bets[k] for k in 1:10))[2]-1

predictions = [classify(testData[k,:]) for k in 1:nTest]
confusionMatrix = [sum((predictions == i) .& (testLabels == j))
for i in 0:9, j in 0:9]
accuracy = 100*sum(diag(confusionMatrix))/nTest

println("Accuracy: ", accuracy,"%")
confusionMatrix
```

**Basic statistics  
(least squares/regression)**



Accuracy: 86.03%

10×10 Array{Int64,2}:

944	0	18	4	0	23	18	5	14	15
0	1107	54	17	22	18	10	40	46	11
1	2	813	23	6	3	9	16	11	2
2	2	26	880	1	72	0	6	30	17
2	3	15	5	881	24	22	26	27	80
7	1	0	17	5	659	17	0	40	1
14	5	42	9	10	23	875	1	15	1
2	1	22	21	2	14	0	884	12	77
7	14	37	22	11	39	7	0	759	4
1	0	5	12	44	17	0	50	20	801

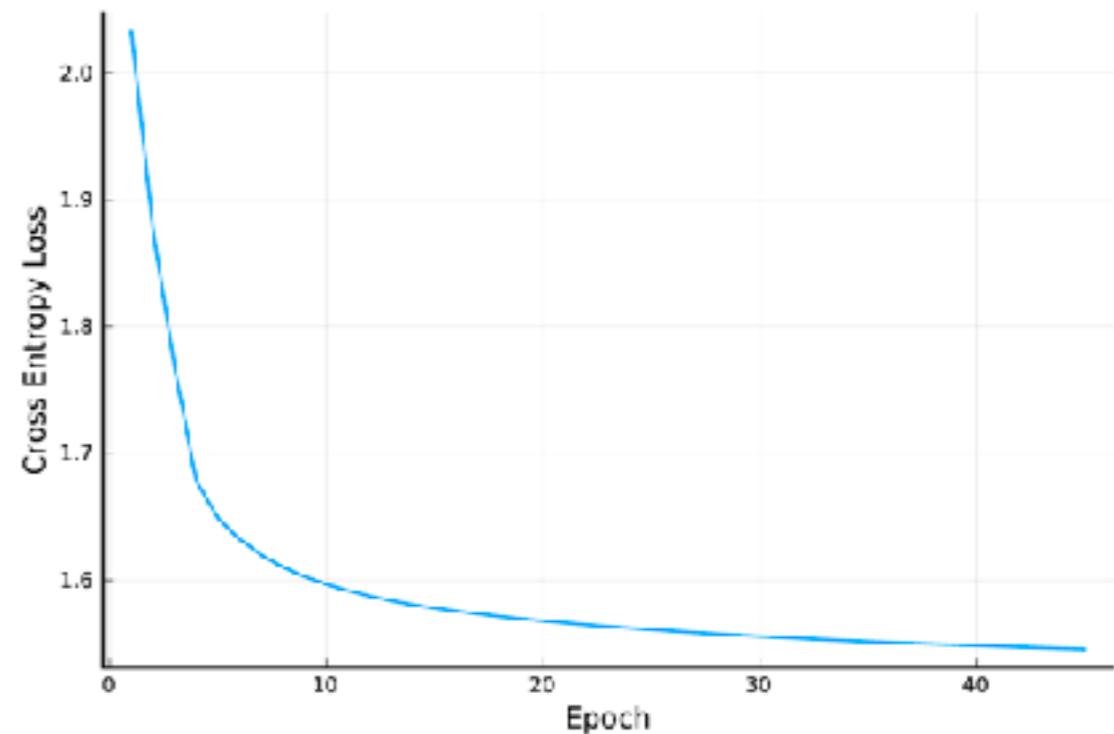
# Logistic softmax regression

$$\sigma(u) = \frac{1}{1 + e^{-u}} = \frac{e^u}{e^u + 1}$$

$$\hat{y}(\tilde{x}) = \operatorname{argmax}_{\ell=0,\dots,9} \sigma(w^{(\ell)} \cdot \tilde{x} + b^{(\ell)}).$$

$$s(z) = \frac{1}{\sum_{j=1}^K e^{z_j}} \begin{bmatrix} e^{z_1} & e^{z_2} & \dots & e^{z_K} \end{bmatrix}^T$$

$$\hat{y}(\tilde{x}) = \operatorname{argmax}_{\ell=0,\dots,9} s_\ell(W\tilde{x} + b)$$



Cross Entropy Loss:  $-\sum_{i=1}^n \log(\hat{y}_{y_i+1})$

# **Part 4: A walk in the random forest**

# Random forest is a very “generic” algorithm

```
: using Flux.Data.MNIST, DecisionTree, Random
Random.seed!(0)

trainImgs    = MNIST.images()
trainLabels  = MNIST.labels()
nTrain = length(trainImgs)
trainData = vcat([hcat(float.(trainImgs[i])...) for i in 1:nTrain]...)

testImgs = MNIST.images(:test)
testLabels = MNIST.labels(:test)
nTest = length(testImgs)
testData = vcat([hcat(float.(testImgs[i])...) for i in 1:nTest]...)

numFeaturesPerTree = 10
numTrees = 40
portionSamplesPerTree = 0.7
maxTreeDepth = 10

model = build_forest(trainLabels, trainData,
                     numFeaturesPerTree, numTrees,
                     portionSamplesPerTree, maxTreeDepth)
println("Trained model:")
println(model)

predicted_labels = [apply_forest(model, testData[k,:]) for k in 1:nTest]
accuracy = sum(predicted_labels .== testLabels)/nTest
println("\nPrediction accuracy (measured on test set of size $nTest): ",accuracy)
```

Trained model:  
Ensemble of Decision Trees  
Trees: 40  
Avg Leaves: 860.525  
Avg Depth: 10.0

Prediction accuracy (measured on test set of size 10000): 0.9377

# **Activity B: MNIST digit classification with random forests**

# **Part 5: Going Deep**

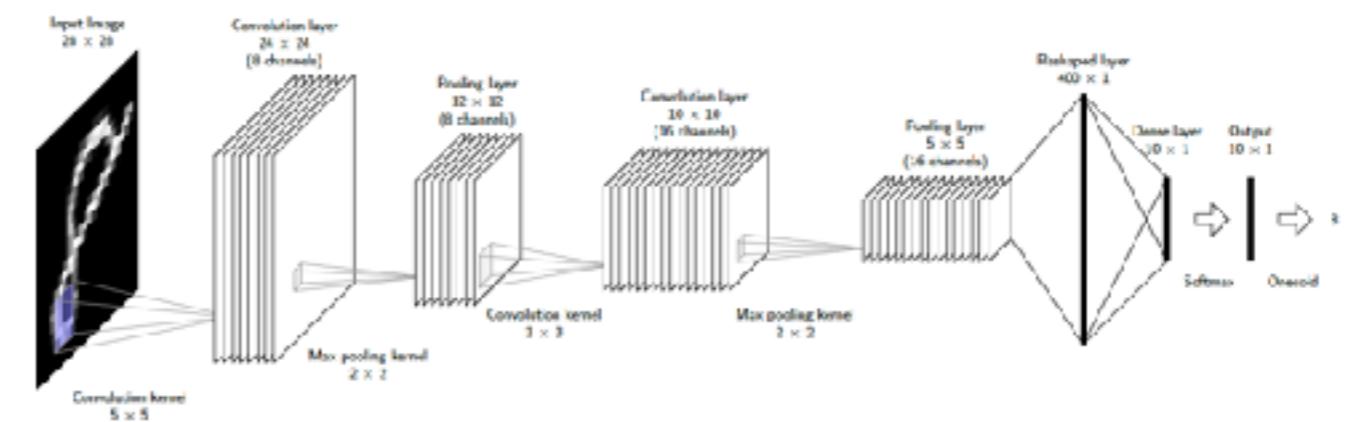
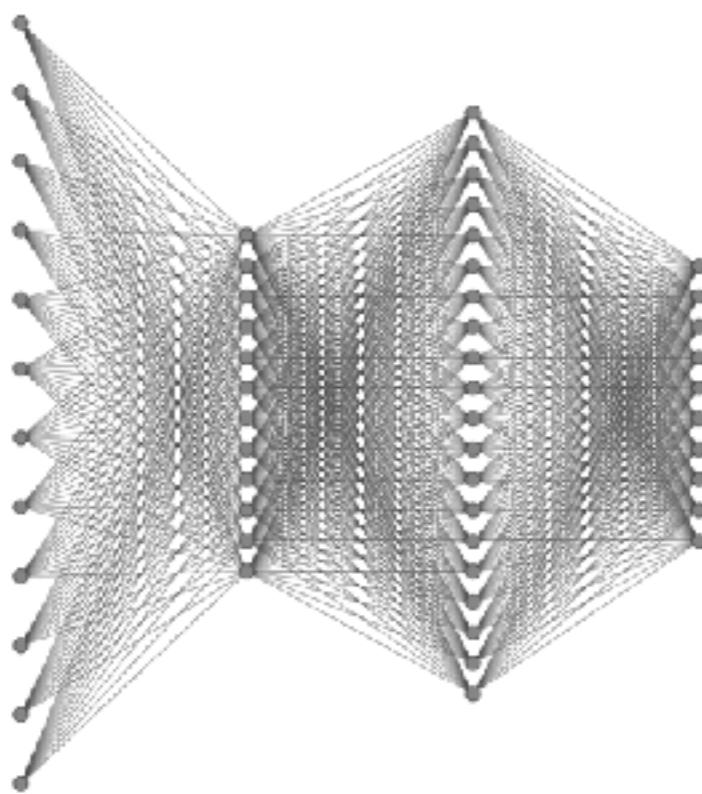
# A Deep Neural Network is just a bit more than a linear model...

$\rho_i(\cdot) \equiv$  Activation function

$$\tilde{f}(x) = s\left( A_3 \underbrace{\rho_2\left(A_2 \underbrace{\rho_1\left(A_1 \overbrace{x}^{\text{Input layer}} + b_1\right) + b_2\right)}_{\text{First hidden layer}} + b_3\right).$$

$\underbrace{\hspace{10em}}$   
 $\underbrace{\hspace{10em}}$   
 $\underbrace{\hspace{10em}}$

Input layer      First hidden layer      Second hidden layer      Output layer



# Training two neural network variants

```
using Flux, Flux.Data.MNIST, Statistics, Random, Plots
using Flux: onehotbatch, onecold, crossentropy, flatten
Random.seed!(0)

epochs = 30
eta = 5e-3
batchSize = 1000
trainRange, validateRange = 1:5000, 5001:10000

function minibatch(x, y, indexRange)
    xBatch = Array{Float32}(undef, size(x[1])..., 1, length(indexRange))
    for i in 1:length(indexRange)
        xBatch[:, :, :, i] = Float32.(x[indexRange[i]])
    end
    return (xBatch, onehotbatch(y[indexRange], 0:9))
end

trainLabels = MNIST.labels()[trainRange]
trainImage = MNIST.images()[trainRange]
mbIdxs = Iterators.partition(1:length(trainImage), batchSize)
trainSet = [minibatch(trainImage, trainLabels, bi) for bi in mbIdxs]

validateLabels = MNIST.labels()[validateRange]
validateImage = MNIST.images()[validateRange]
validateSet = minibatch(validateImage, validateLabels, 1:length(validateImage))

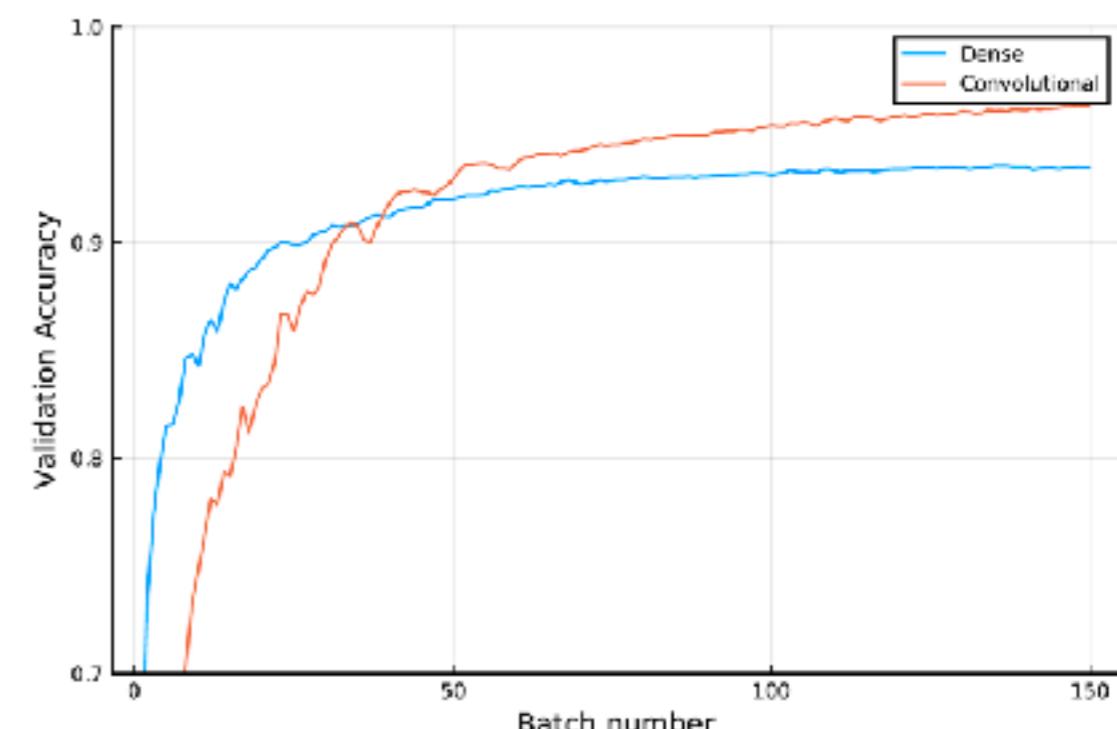
model1 = Chain(flatten, Dense(784, 200, relu), Dense(200, 100, tanh),
               Dense(100, 10, sigmoid), softmax)

model2 = Chain(Conv((5, 5), 1=>8, relu), MaxPool((2,2)),
               Conv((3, 3), 8=>16, relu), MaxPool((2,2)),
               flatten, Dense(400, 10), softmax)

opt1 = ADAM(eta); opt2 = ADAM(eta)
accuracyPaths = [[], []]
accuracy(x, y, model) = mean(onecold(model(x)) .== onecold(y))
loss(x, y, model) = crossentropy(model(x), y)
cbF1() = push!(accuracyPaths[1], accuracy(validateSet..., model1))
cbF2() = push!(accuracyPaths[2], accuracy(validateSet..., model2))

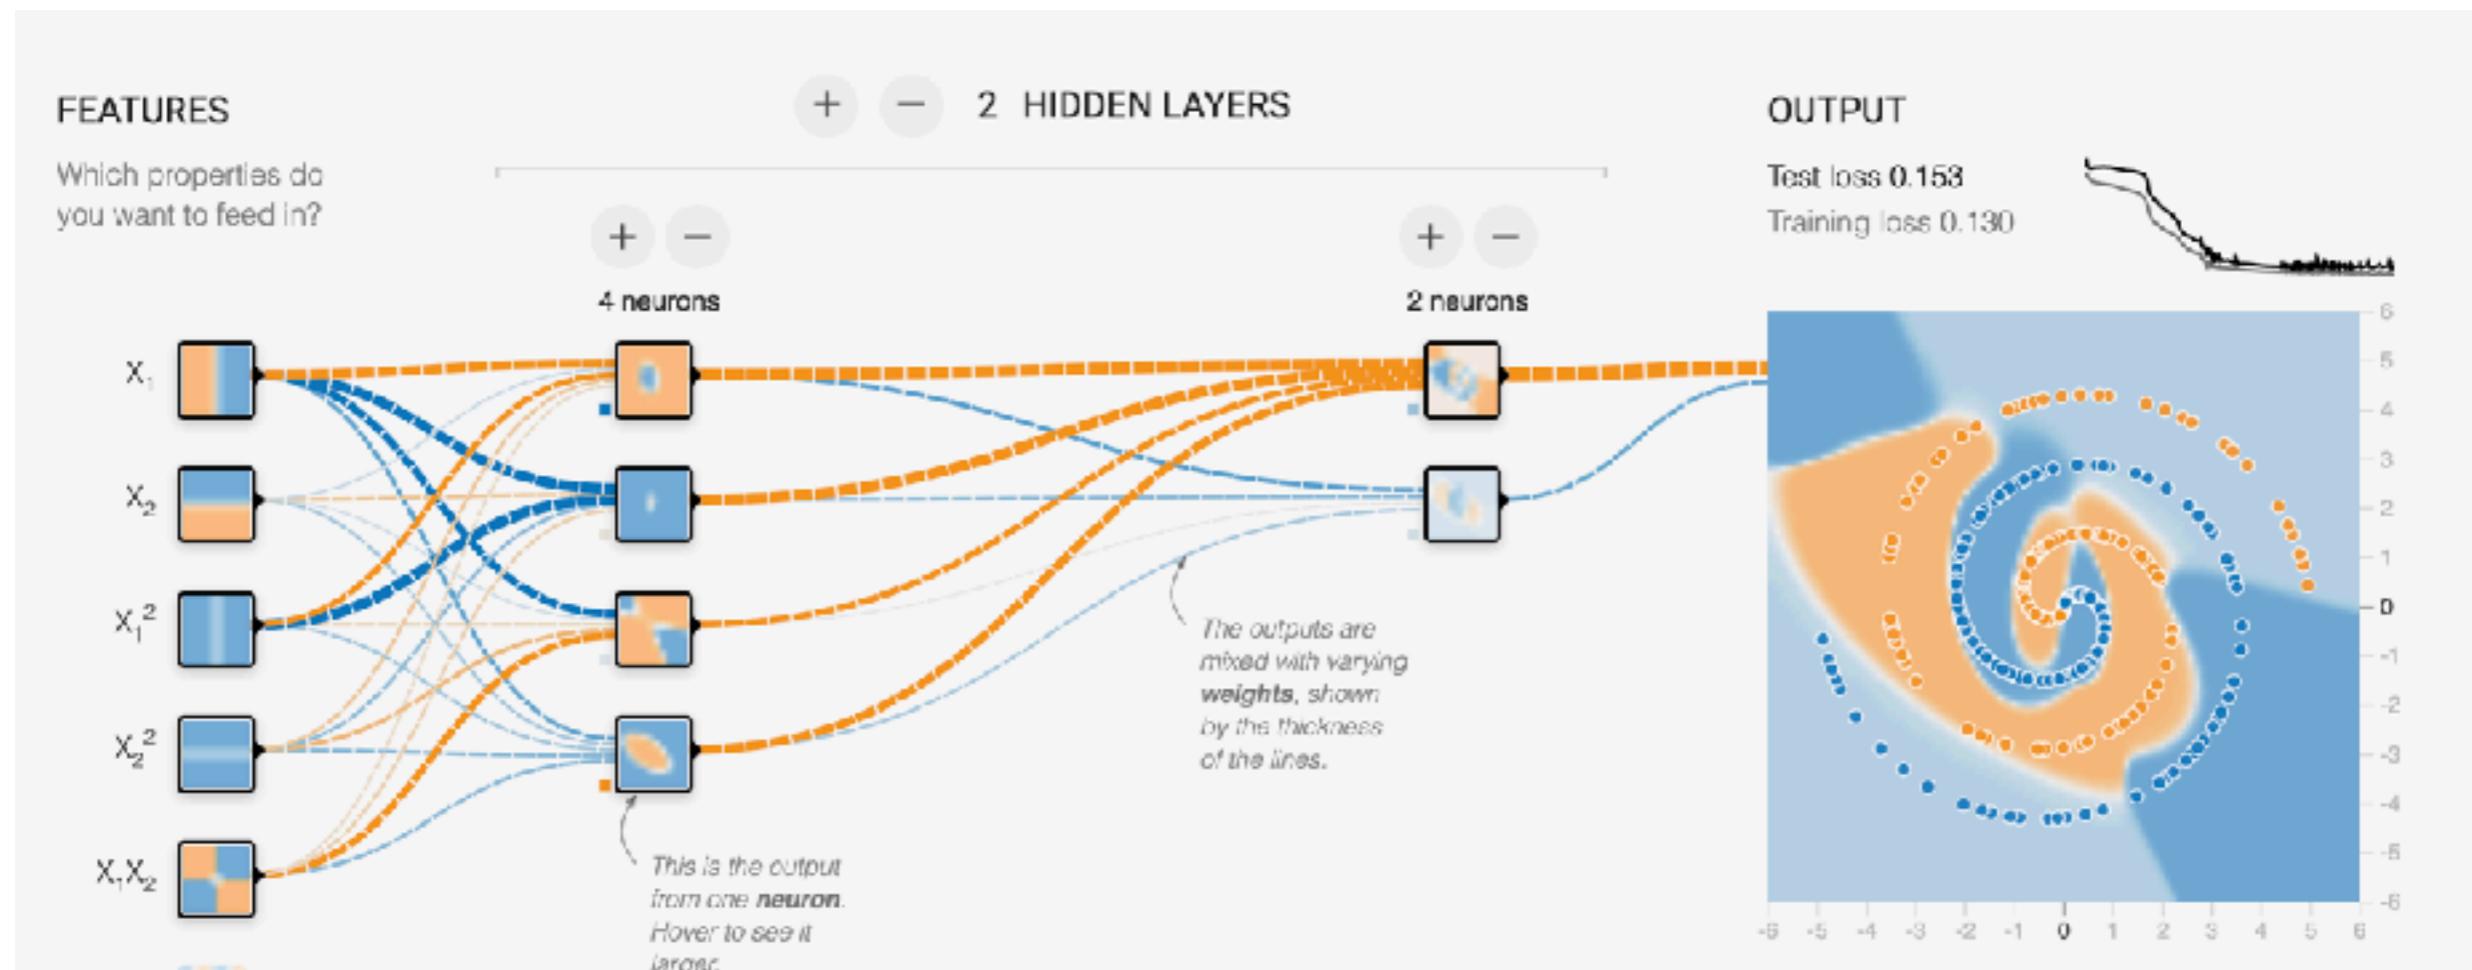
model1!(trainSet[1][1]); model2!(trainSet[1][1])
for _ in 1:epochs
    Flux.train!((x,y)=>loss(x,y,model1), params(model1), trainSet, opt1, cb=cbF1)
    Flux.train!((x,y)=>loss(x,y,model2), params(model2), trainSet, opt2, cb=cbF2)
    print(".")
end

println("\nModel1 (Dense) accuracy = ", accuracy(validateSet..., model1))
println("Model2 (Convolutional) accuracy = ", accuracy(validateSet..., model2))
plot(accuracyPaths, label = ["Dense" "Convolutional"],
      ylim=(0.7,1.0), xlabel="Batch number", ylabel = "Validation Accuracy")
```



# Activity C: Google Tensor flow Playground

<https://playground.tensorflow.org/>



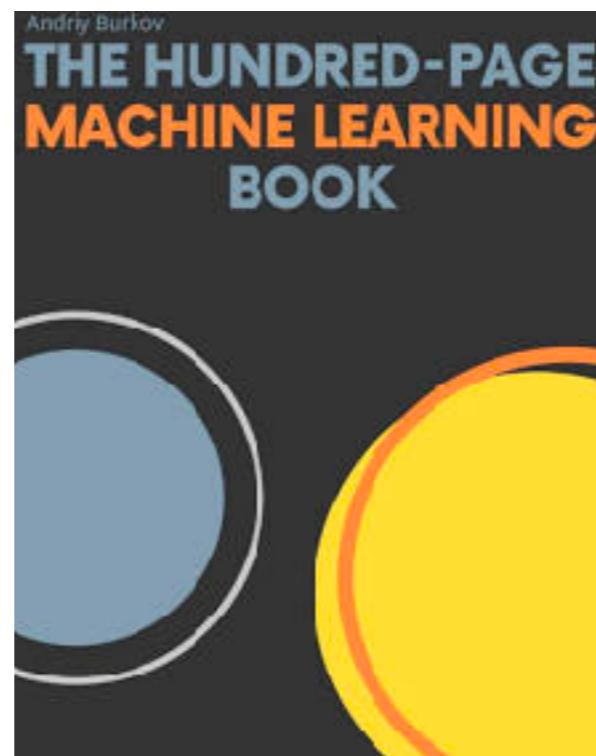
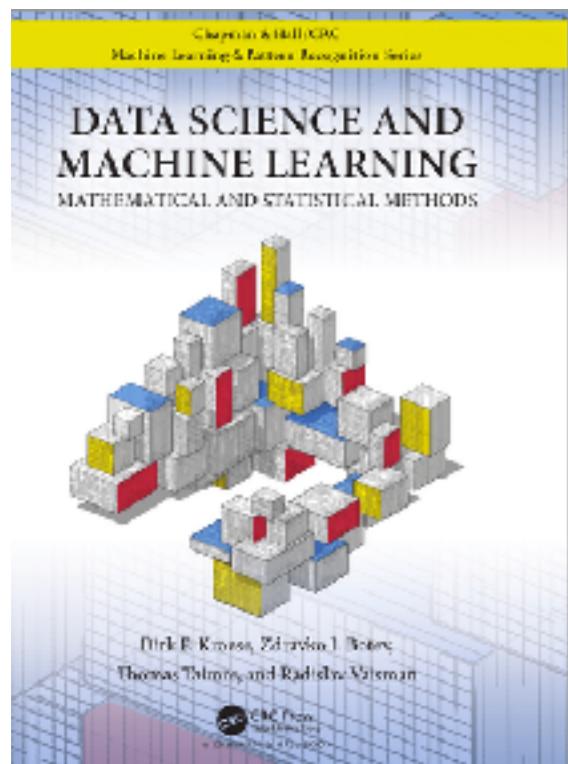
# **Activity D: MNIST Digit Classification with dense NN**

# **Activity E: MNIST Digit Classification with conv NN**

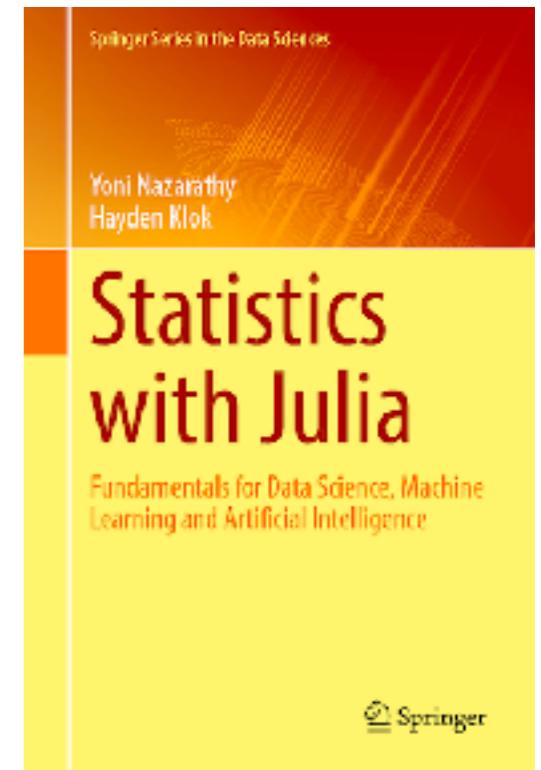
# **Activity F: CIFAR10 Digit Classification with conv NN**

# Closing

# Some resources



**The Mathematical  
Engineering of Deep  
Learning**  
**Benoit Liquet, Sarat  
Moka, and Yoni  
Nazarathy**



<https://people.smp.uq.edu.au/DirkKroese/DSML/>

<http://themlbook.com/>

<https://deeplearningmath.org/>

<https://statisticswithjulia.org/>